

---

# Transport Overview

KR 3.1-3.4

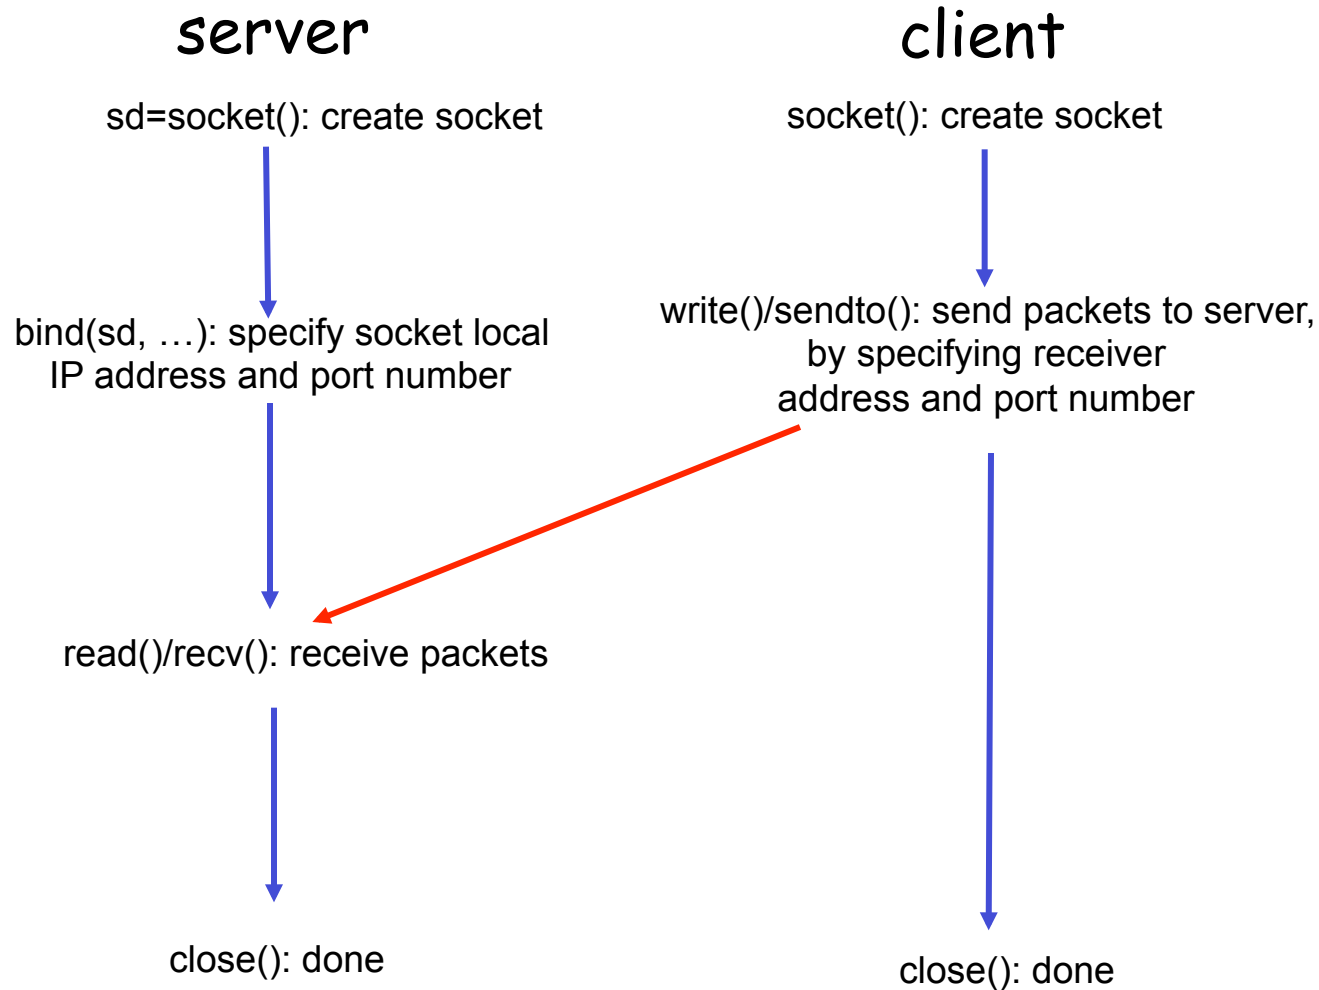
# Outline

---

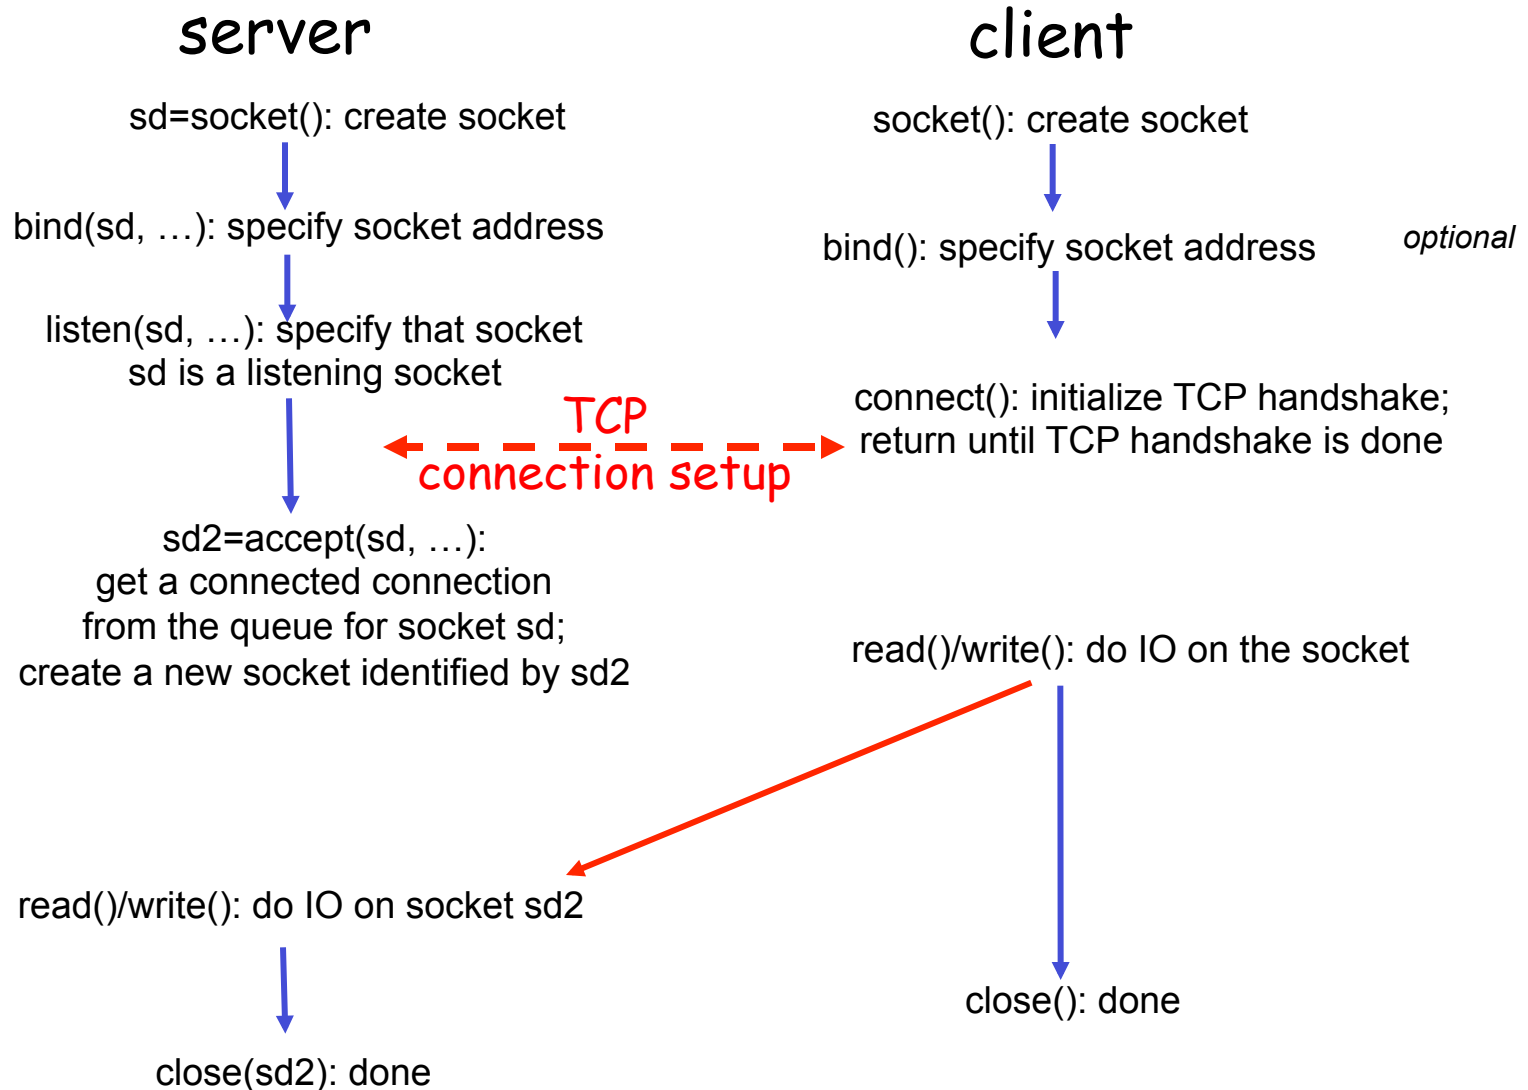
## ➤ Recap

- ❑ Overview of transport layer
- ❑ UDP
- ❑ Reliable data transfer, the stop-and-wait protocol

# Connectionless UDP: Big Picture



# Connection-oriented: Big Picture



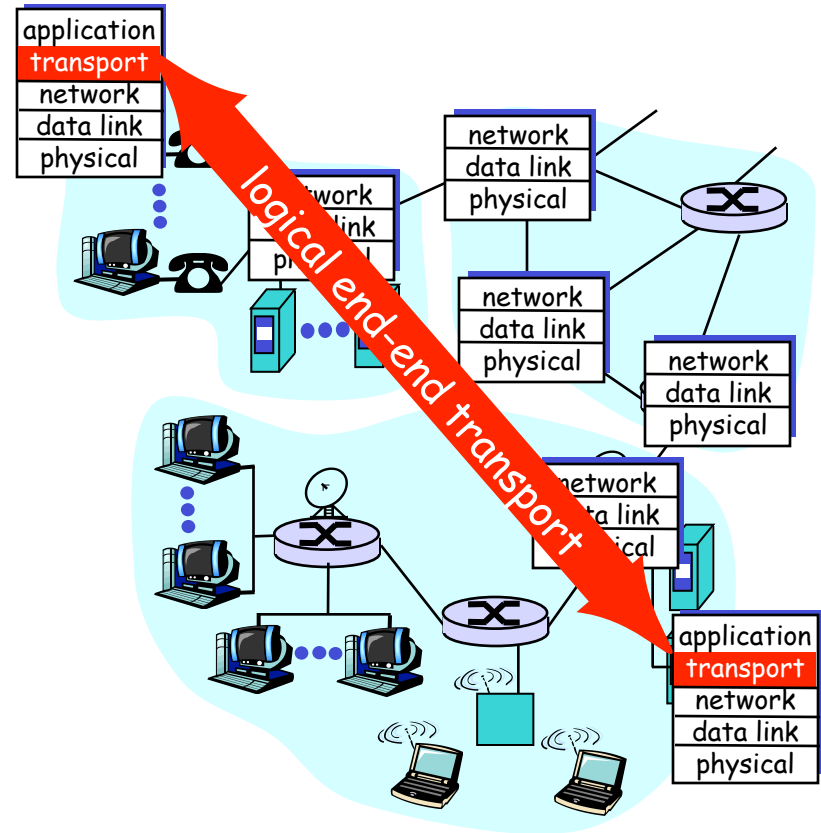
# Outline

---

- ❑ Recap
  - Overview of transport layer
- ❑ UDP
- ❑ Reliable data transfer, the stop-and-go protocols

# Transport Layer vs. Network Layer

- ❑ Provide *logical communication* between app' processes
- ❑ Transport protocols run in end systems
  - Sender: breaks app messages into **segments**, passes to network layer
  - Receiver: reassembles segments into messages, passes to app layer
- ❑ **Transport vs. network layer services:**
  - *Network layer*: data transfer between end systems
  - *Transport layer*: data transfer between processes
    - Relies on, enhances network layer services



# Transport Layer Services and Protocols

---

- ❑ Reliable, in-order delivery (TCP)
  - Multiplexing
  - Reliability and connection setup
  - Congestion control
  - Flow control
  
- ❑ Unreliable, unordered delivery: UDP
  - Multiplexing
  
- ❑ Services not available:
  - Delay guarantees
  - Bandwidth guarantees

# Transport Layer: Road Ahead

- ❑ Class 1:
  - Connectionless transport: UDP
  - Reliable data transfer using stop-and-wait
- ❑ Class 2:
  - Sliding window reliability
  - TCP reliability
    - Overview of TCP
    - TCP RTT measurement
    - TCP connection management
- ❑ Class 3:
  - Principles of congestion control
  - TCP congestion control; AIMD
- ❑ Class 4:
  - Performance modeling; TCP variants
  - The analysis and design framework for congestion control



# Outline

---

- ❑ Recap
- ❑ Overview of transport layer
  - UDP
- ❑ Reliable data transfer, the stop-and-go protocols

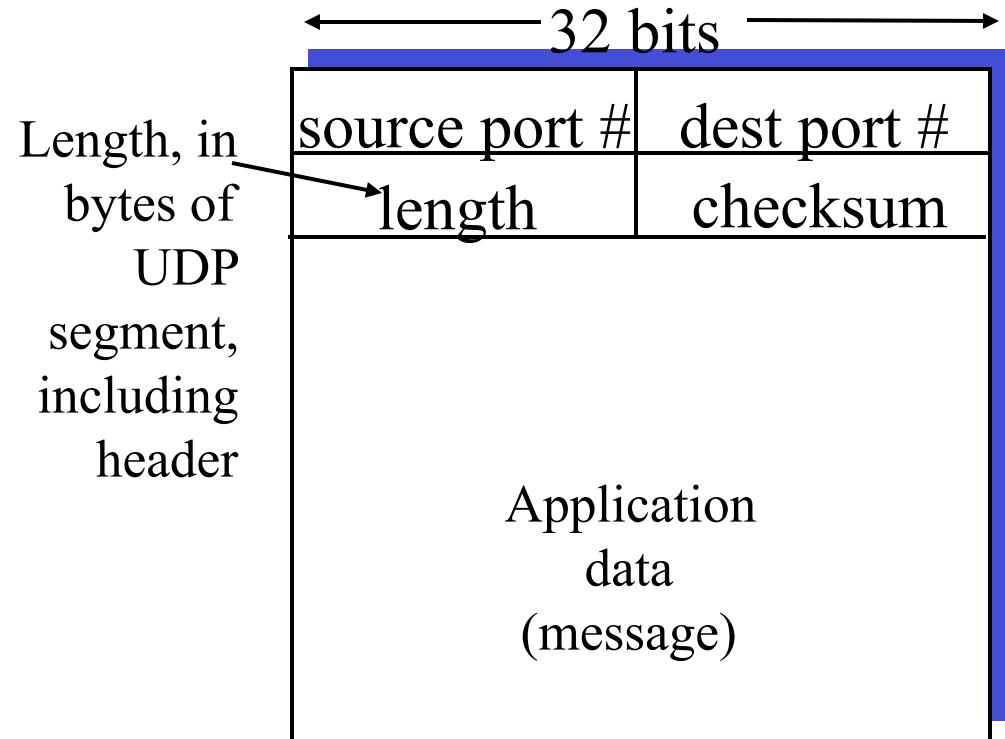
# UDP: User Datagram Protocol [RFC 768]

## ❑ Often used for streaming multimedia apps

- Loss tolerant
- Rate sensitive

## ❑ Other UDP usage

- DNS
- SNMP



UDP segment format

# UDP Checksum

Goal: end-to-end detection of “errors” (e.g., flipped bits) in transmitted segment

## Sender:

- ❑ Treat segment contents as sequence of 16-bit integers (represented in one's complement representation)
- ❑ Checksum: addition (1's complement sum) of segment contents to be  
1111111111111111
- ❑ Sender puts checksum value into UDP checksum field

## Receiver:

- ❑ Compute checksum of received segment
- ❑ Compute sum of segment and checksum; check if sum is 1111111111111111
  - NO - error detected
  - YES - no error detected. *But maybe errors nonetheless?*

# One's Complement Arithmetic

- ❑ UDP checksum is based on one's complement arithmetic
  - One's complement was a common representation of **signed** numbers in early computers
- ❑ One's complement representation
  - Bit-wise NOT for negative numbers
  - Example: assume 8 bits
    - 00000000: 0
    - 00000001: 1
    - 01111111: 127
    - 10000000: ?
    - 11111111: ?
  - Addition: conventional binary addition except adding any resulting carry back into the resulting sum
    - Example:  $-1 + 2$

# UDP Checksum: Algorithm

## □ Example checksum:

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

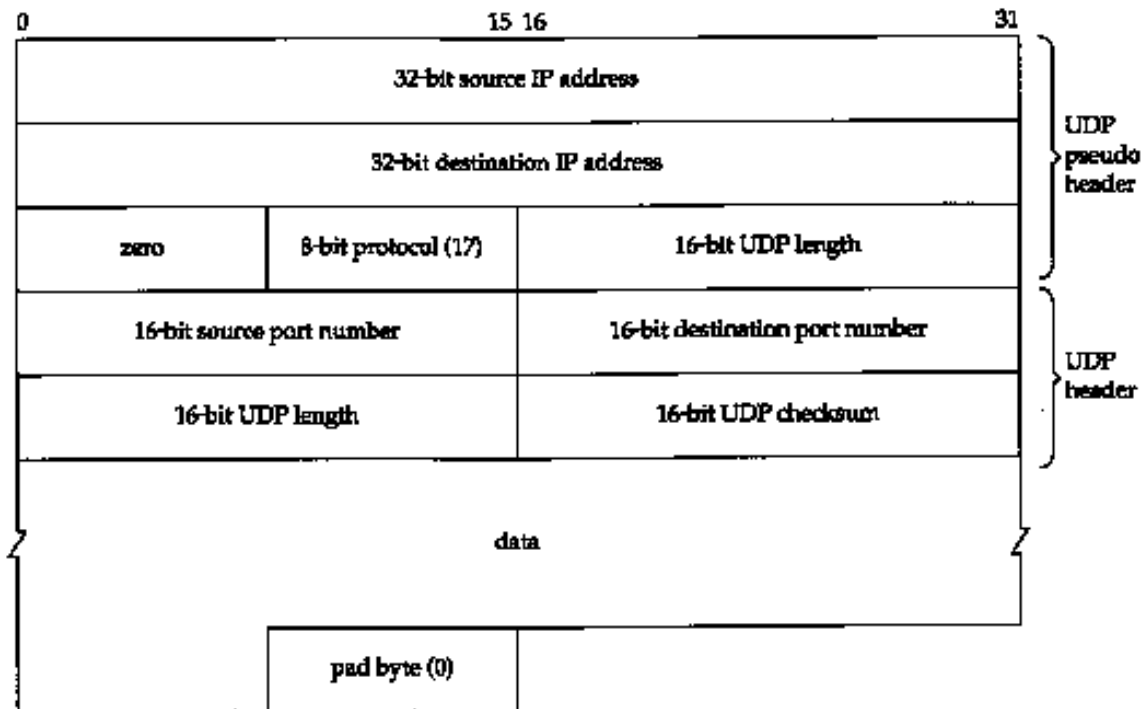
-What are some desirable properties of using one's complement representation?

- For fast implementation of computing UDP checksum, see <http://www.faqs.org/rfcs/rfc1071.html>

# UDP Checksum: Coverage

Calculated over:

- ❑ A pseudo-header
  - IP Source Address (4 bytes)
  - IP Destination Address (4 bytes)
  - Protocol (2 bytes)
  - UDP Length (2 bytes)
- ❑ UDP header
- ❑ UDP data




# Outline

---


- ❑ Recap
- ❑ Overview of transport layer
- ❑ UDP
- Reliable data transfer

# Principles of Reliable Data Transfer

- ❑ Important in app., transport, link layers
- ❑ Top-10 list of important networking topics!



(a) provided service

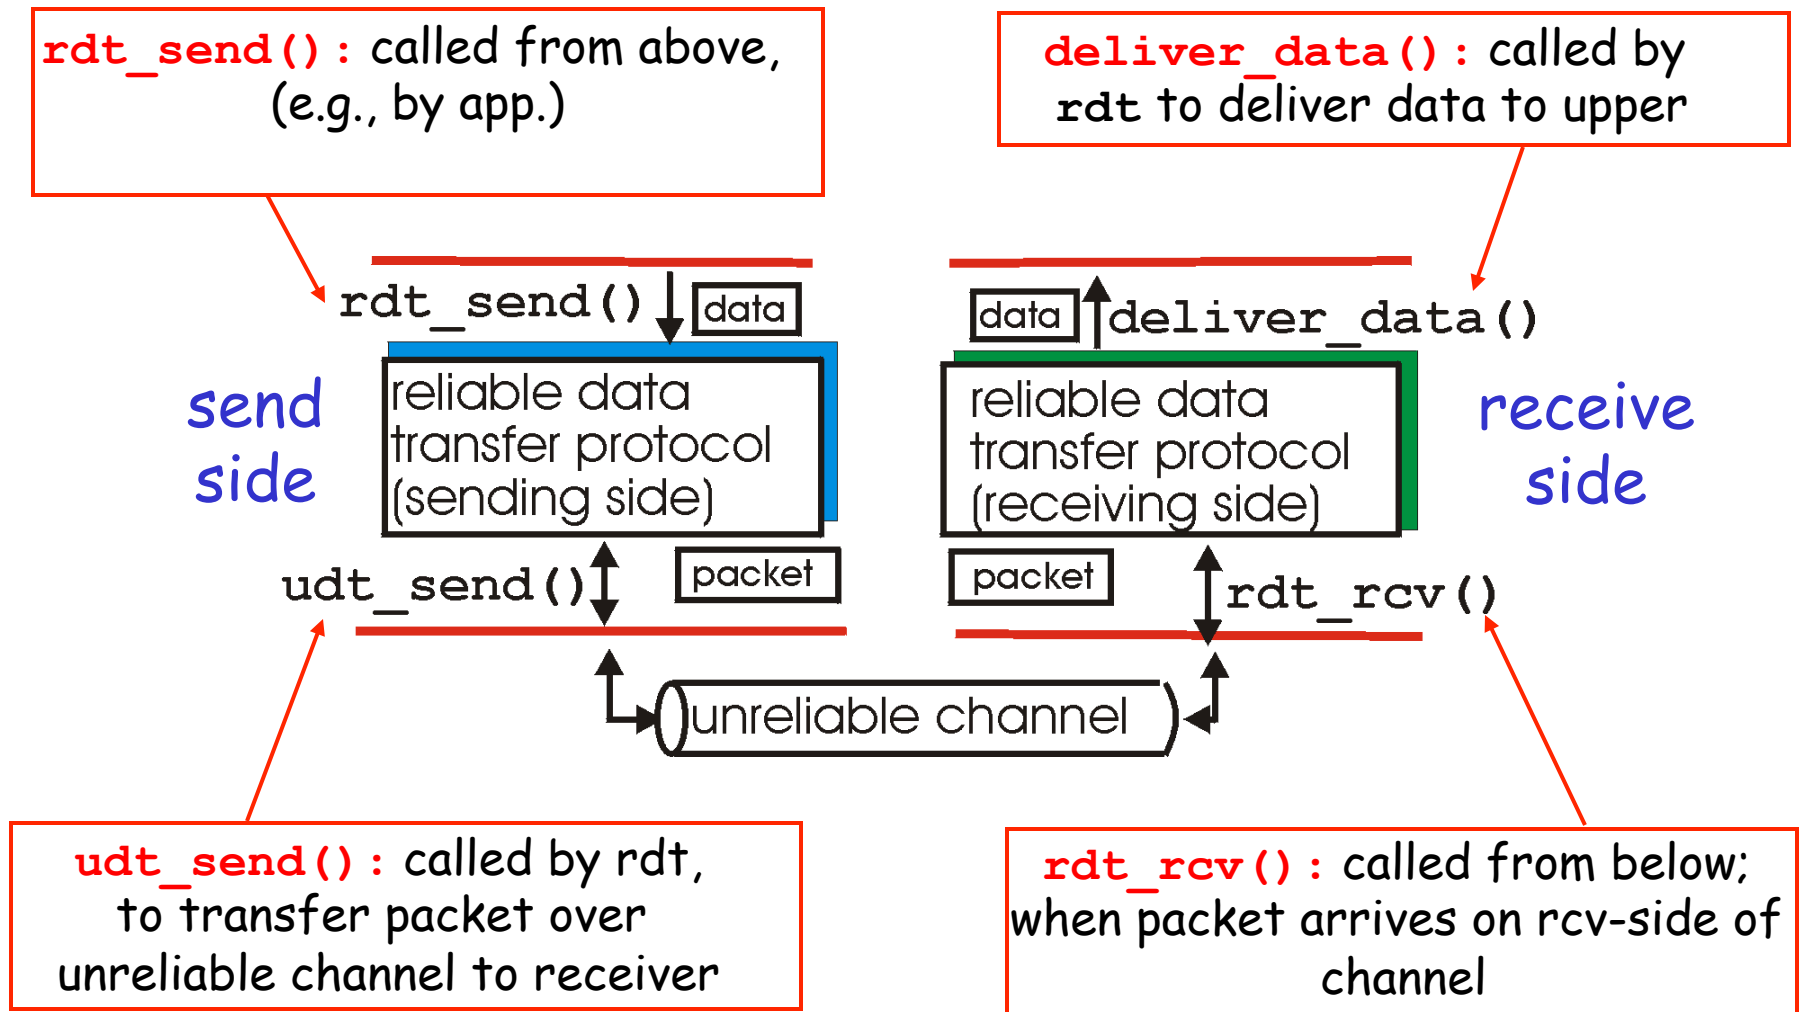


(b) service implementation

Characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt).



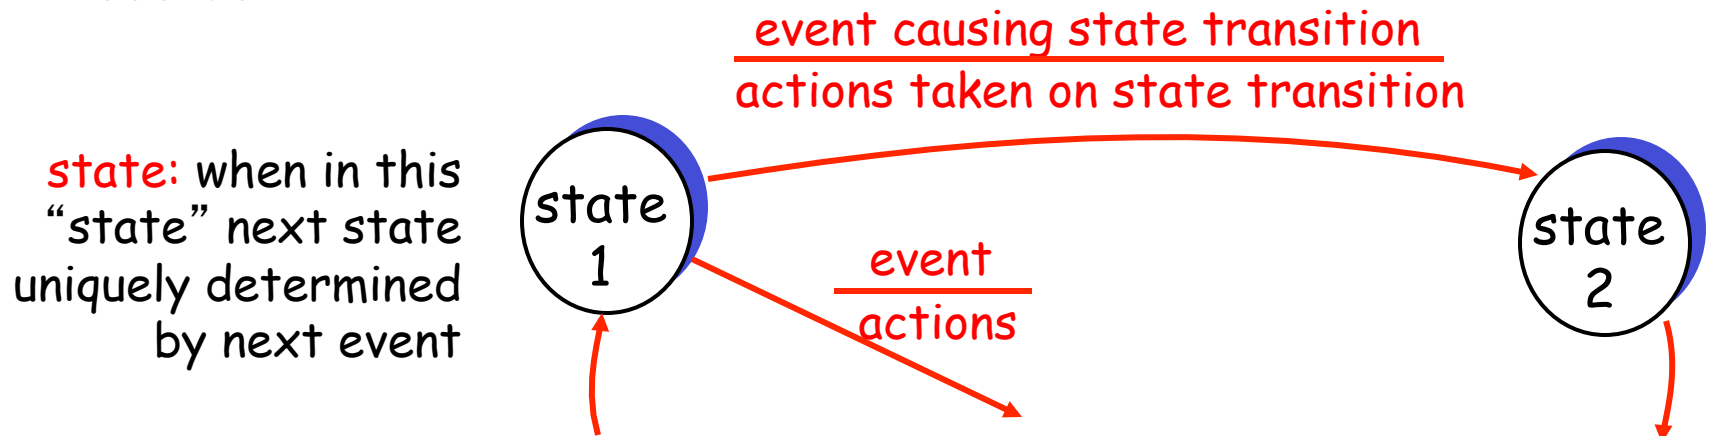
# Reliable Data Transfer: Getting Started



# Reliable Data Transfer: Getting Started

We'll:

- Incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- Consider only unidirectional data transfer
  - But control info will flow on both directions!
- Use finite state machines (FSM) to specify sender, receiver




# Outline

---


- ❑ Review
- ❑ Overview of transport layer
- ❑ UDP
- Reliable data transfer
  - Perfect channel

# Rdt1.0: Reliable Transfer over a Reliable Channel

- Underlying channel perfectly reliable
  - No bit errors
  - No loss of packets
  - No reordering or duplication
- Separate FSMs for sender, receiver:
  - Sender sends data into underlying channel
  - Receiver reads data from underlying channel



sender



receiver

# Outline

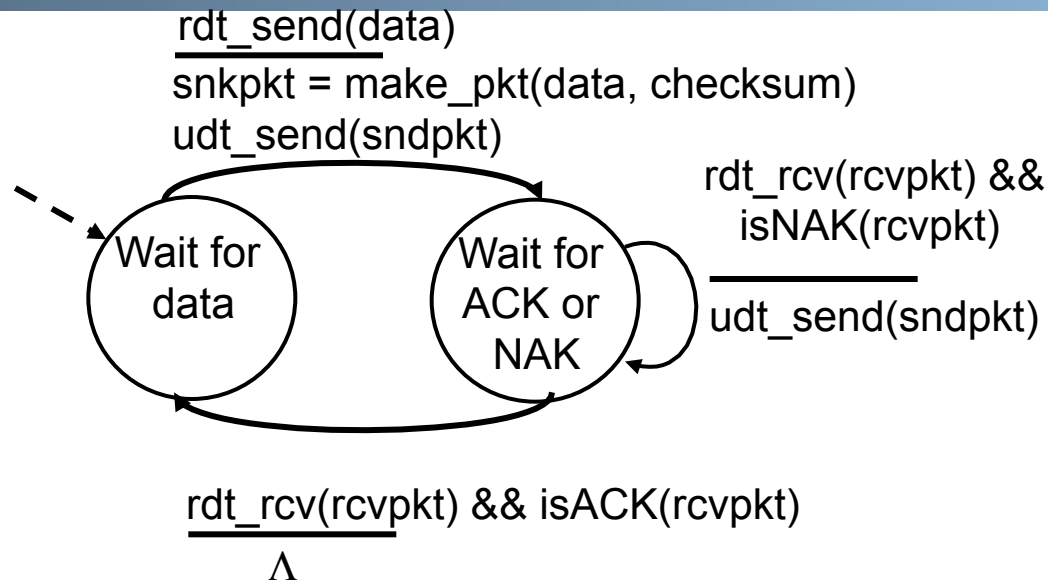
---

- ❑ Recap
- ❑ Overview of transport layer
- ❑ UDP
- Reliable data transfer
  - Perfect channel
  - Channel with bit errors

# Rdt2.0: Channel With Bit Errors


- ❑ Underlying channel may flip bits in packet
  - No loss, duplication or reordering (reasonable?)
- ❑ *The question: how to recover from data pkt errors:*
  - *Acknowledgements (ACKs):* receiver explicitly tells sender that pkt received OK
  - *Negative acknowledgements (NAKs):* receiver explicitly tells sender that pkt had errors
    - Sender retransmits pkt on receipt of NAK
- ❑ New mechanisms in **rdt2.0** (beyond **rdt1.0**):
  - Error detection: recall: UDP checksum to detect bit errors
  - Receiver feedback: control msgs (ACK,NAK) rcvr->sender

# rdt2.0: FSM Specification

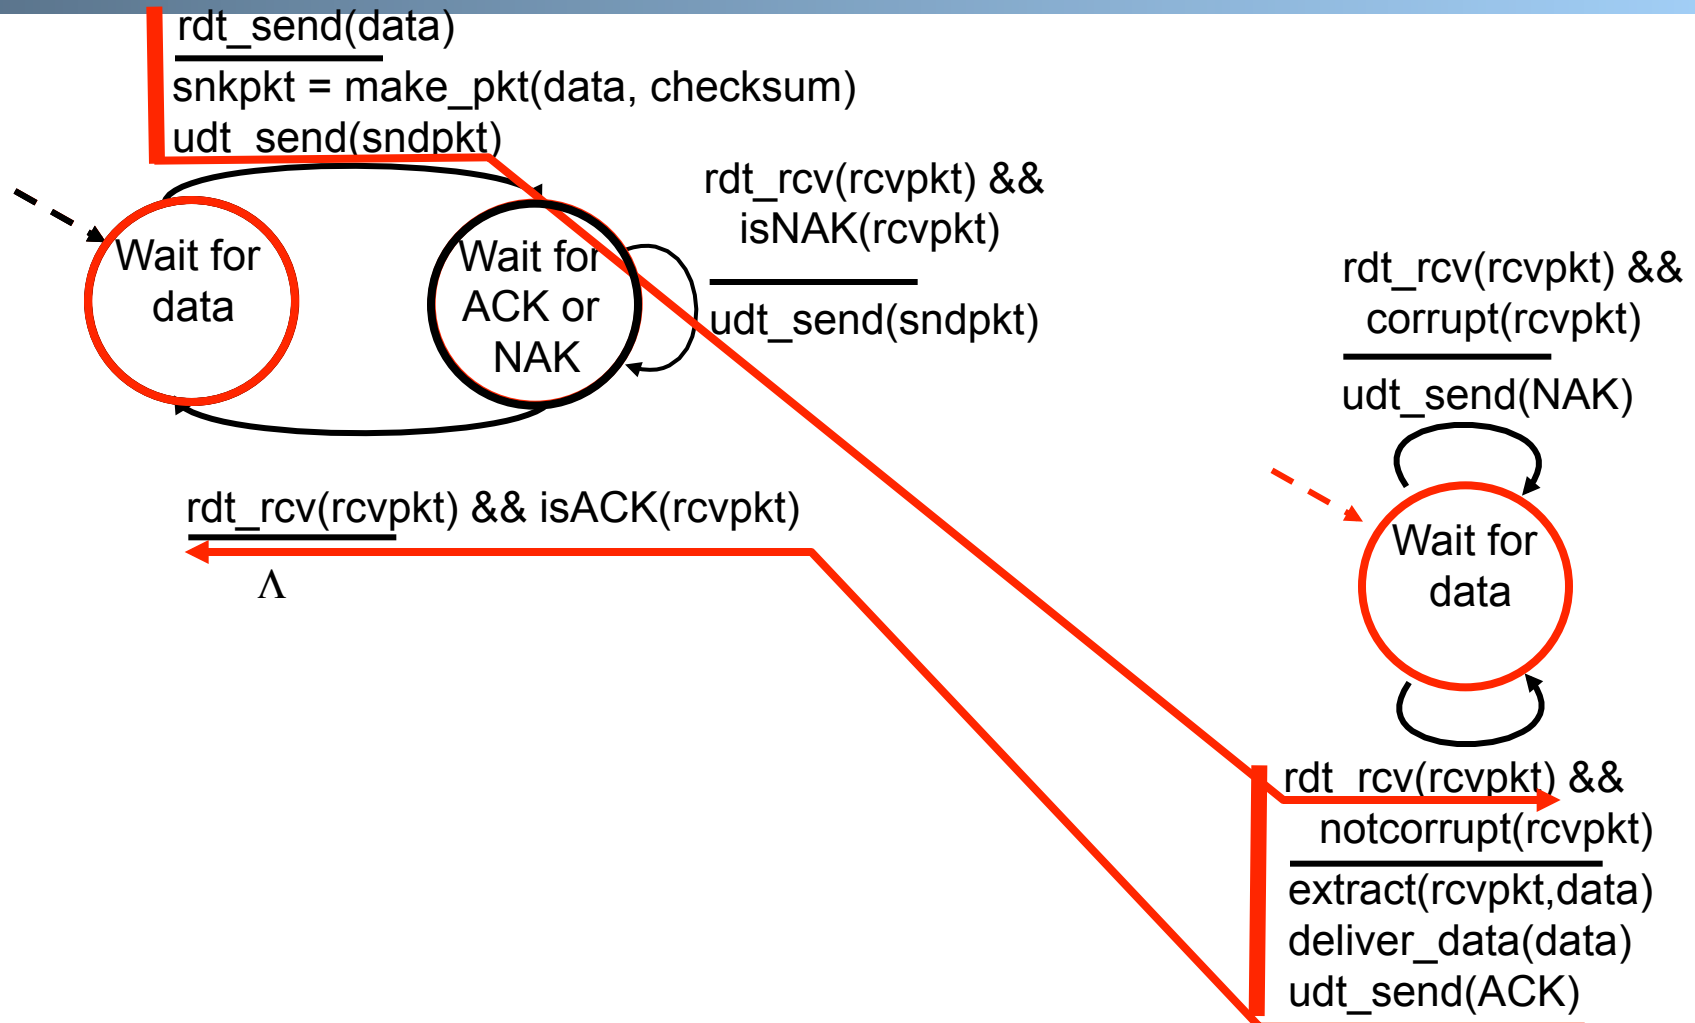


sender

receiver

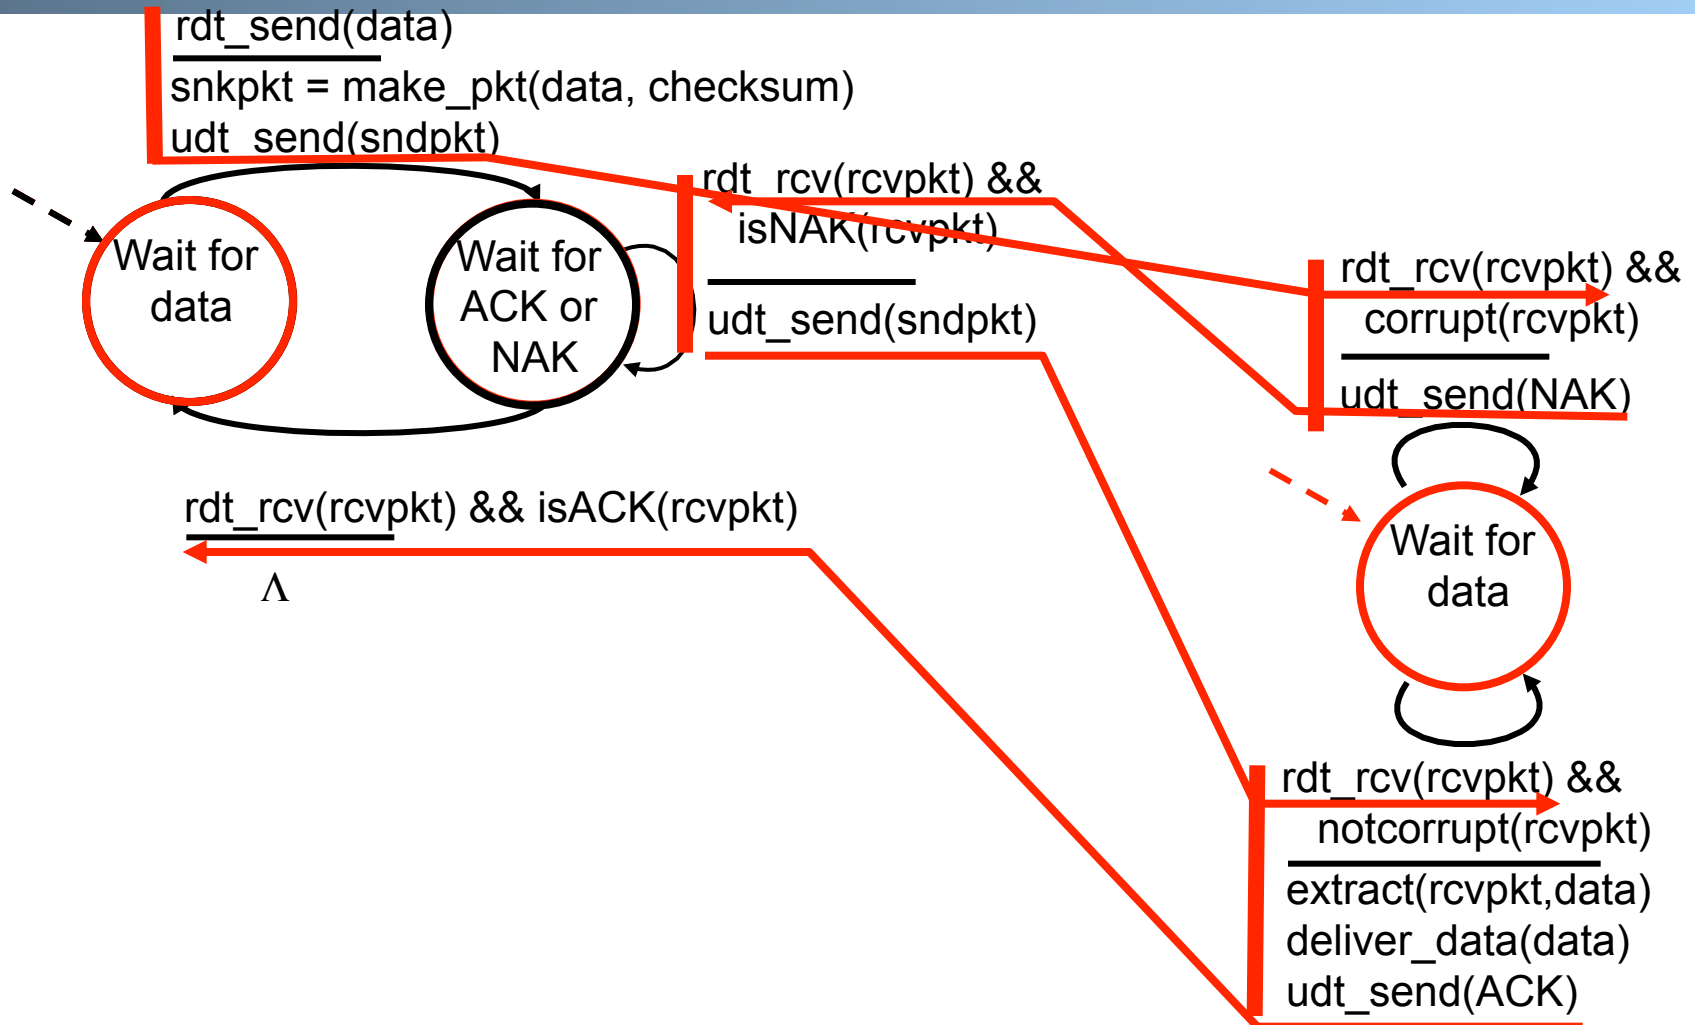


# rdt2.0: Operation with No Errors

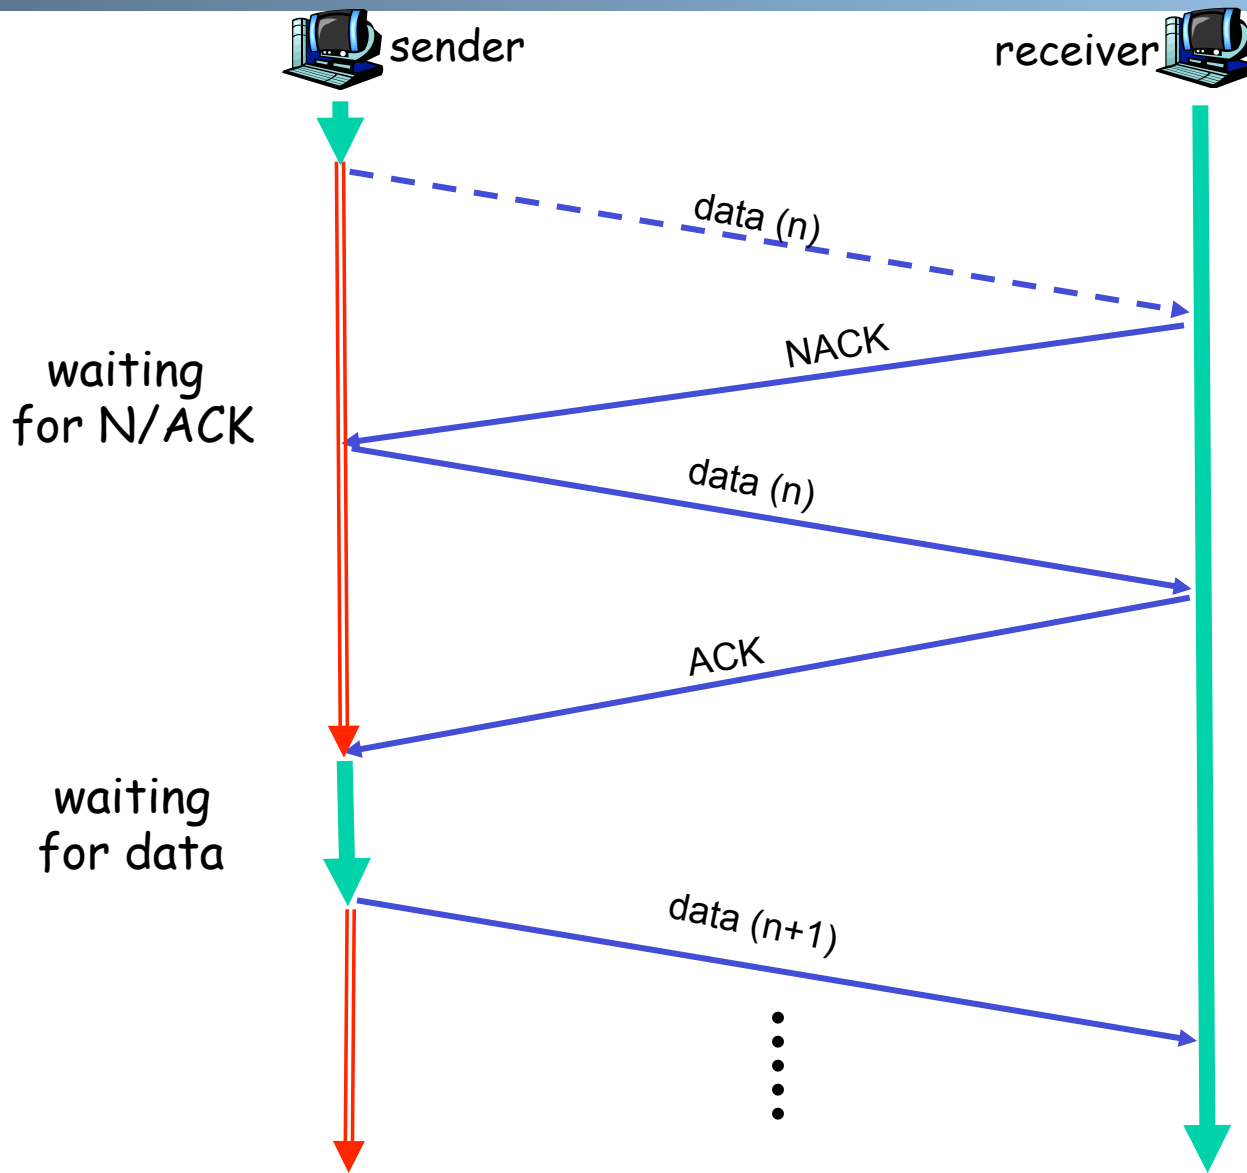




# rdt2.0: Error Scenario




# Big Picture of rdt2.0





# rdt2.0 is Incomplete!

## What happens if ACK/NAK corrupted?

- Although sender receives feedback, but doesn't know what happened at receiver!



# Two Possibilities



# Handle Control Message Corruption

It is always harder to deal with control message errors than data message errors

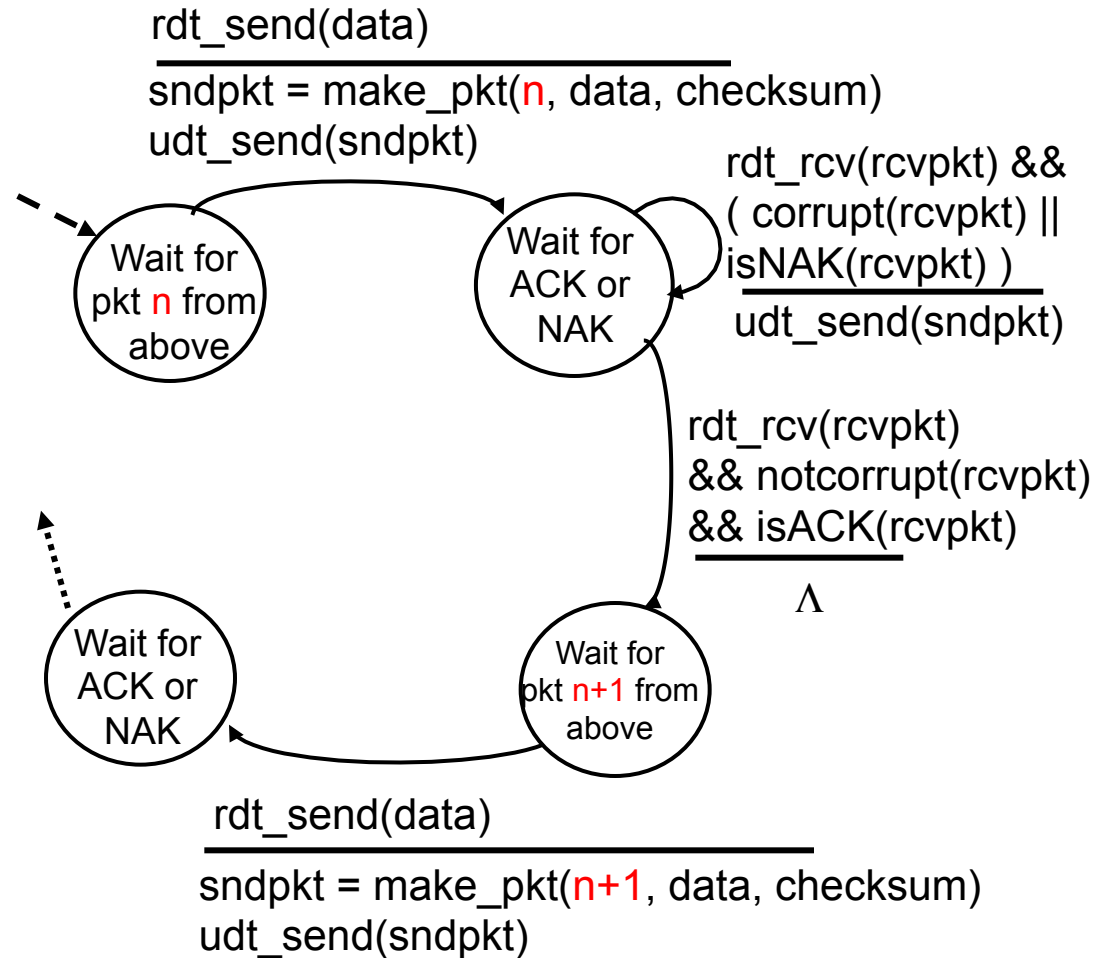
- ❑ Sender can't just retransmit: possible duplicate

## Handling duplicates:

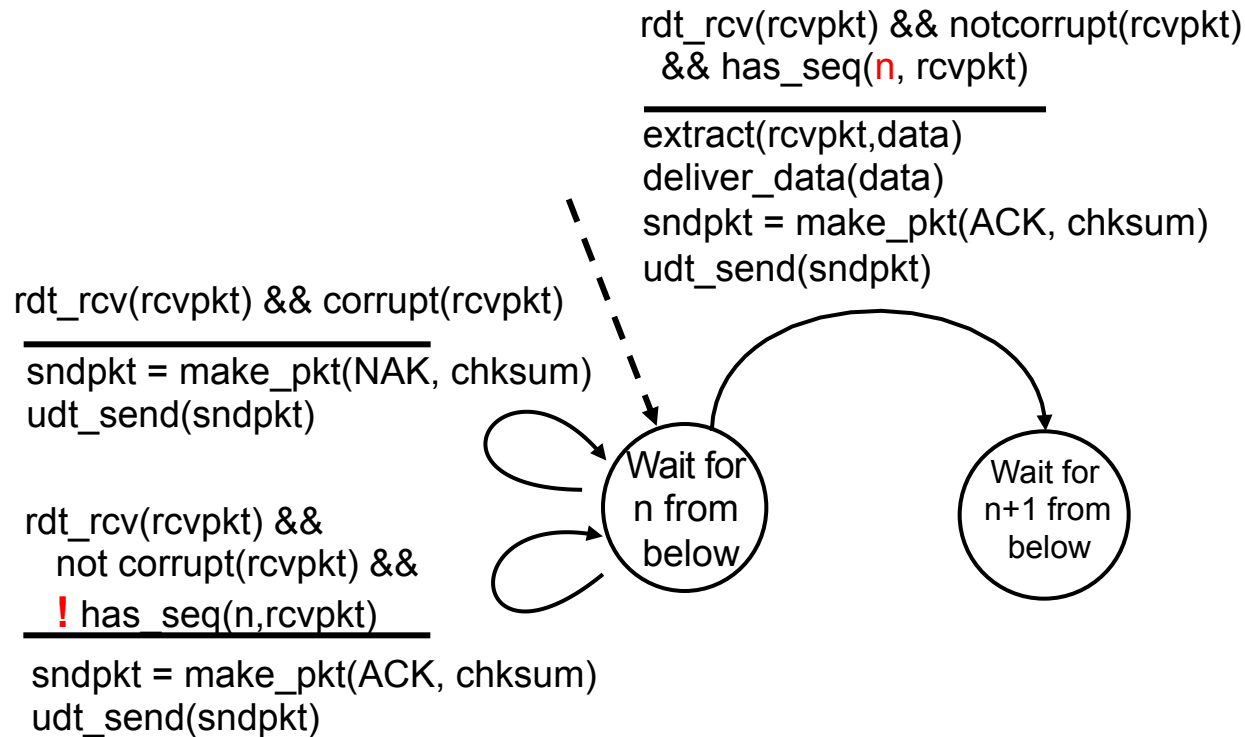
- ❑ Sender adds *sequence number* to each pkt
- ❑ Sender retransmits current pkt if ACK/NAK garbled
- ❑ Receiver discards (doesn't deliver up) duplicate pkt

**stop and wait**  
sender sends one packet,  
then waits for receiver  
response

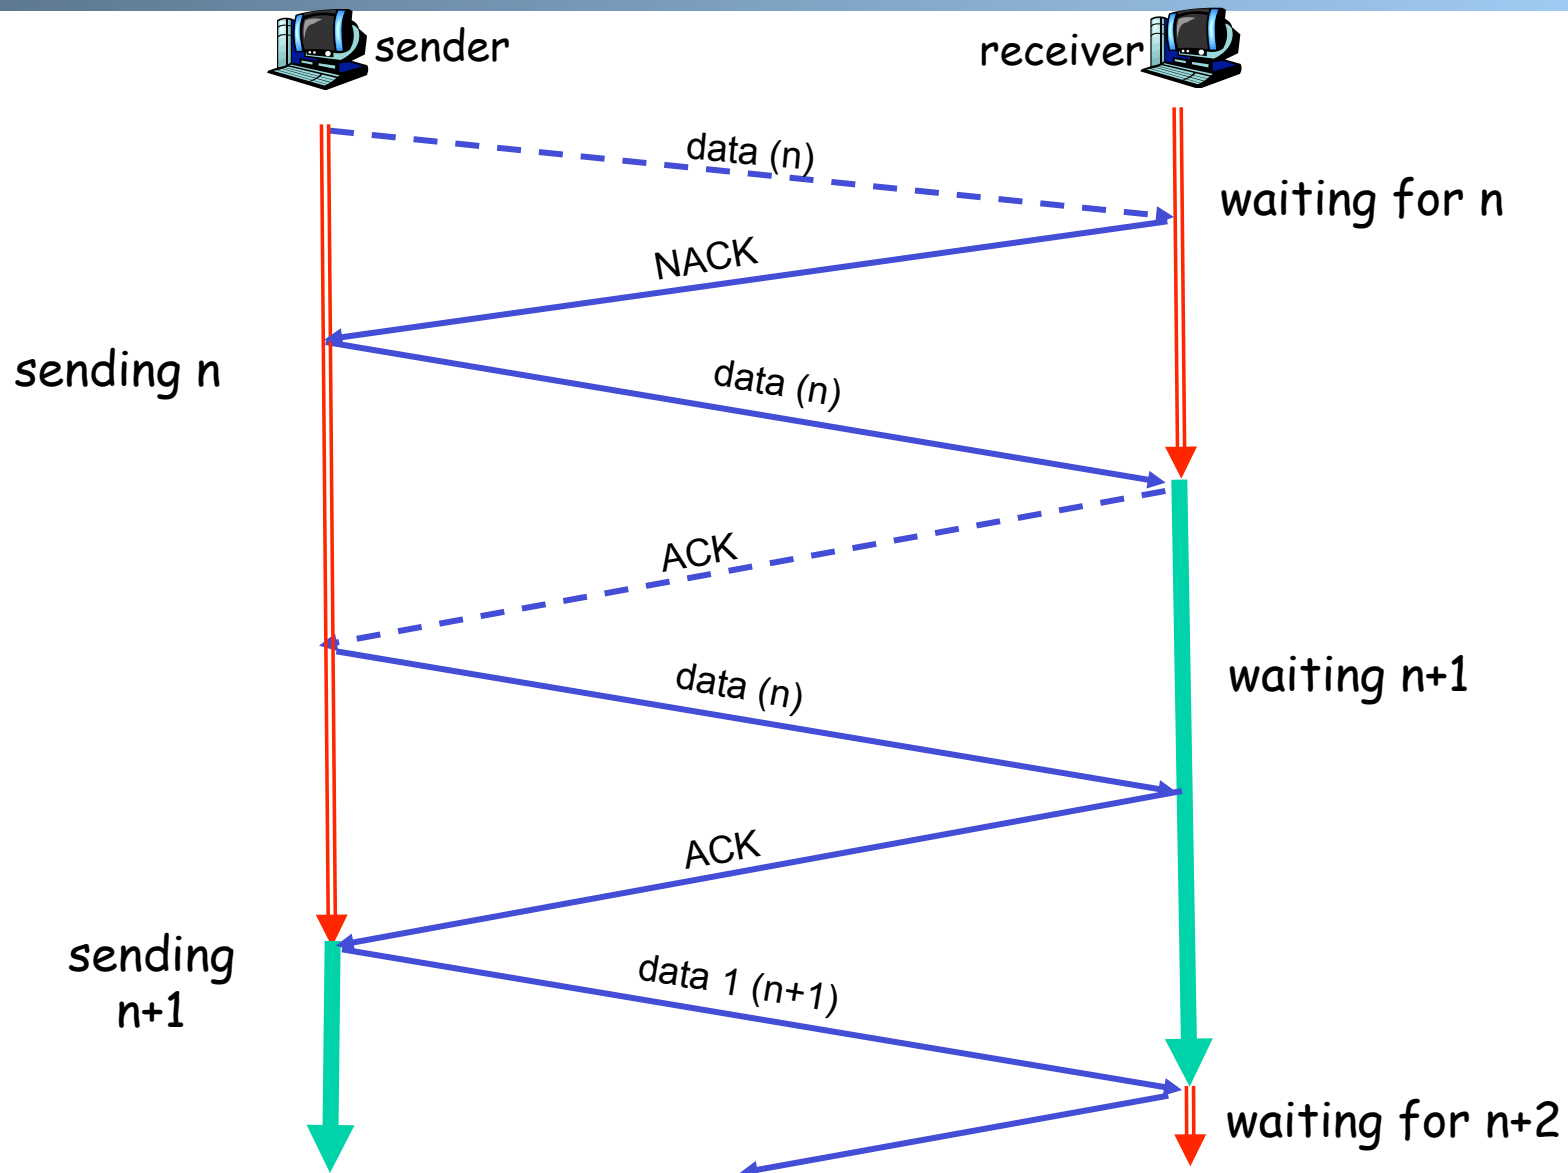
# rdt2.1b: Sender, Handles Garbled ACK/ NAKs



# rdt2.1b: Receiver, Handles Garbled ACK/ NAKs

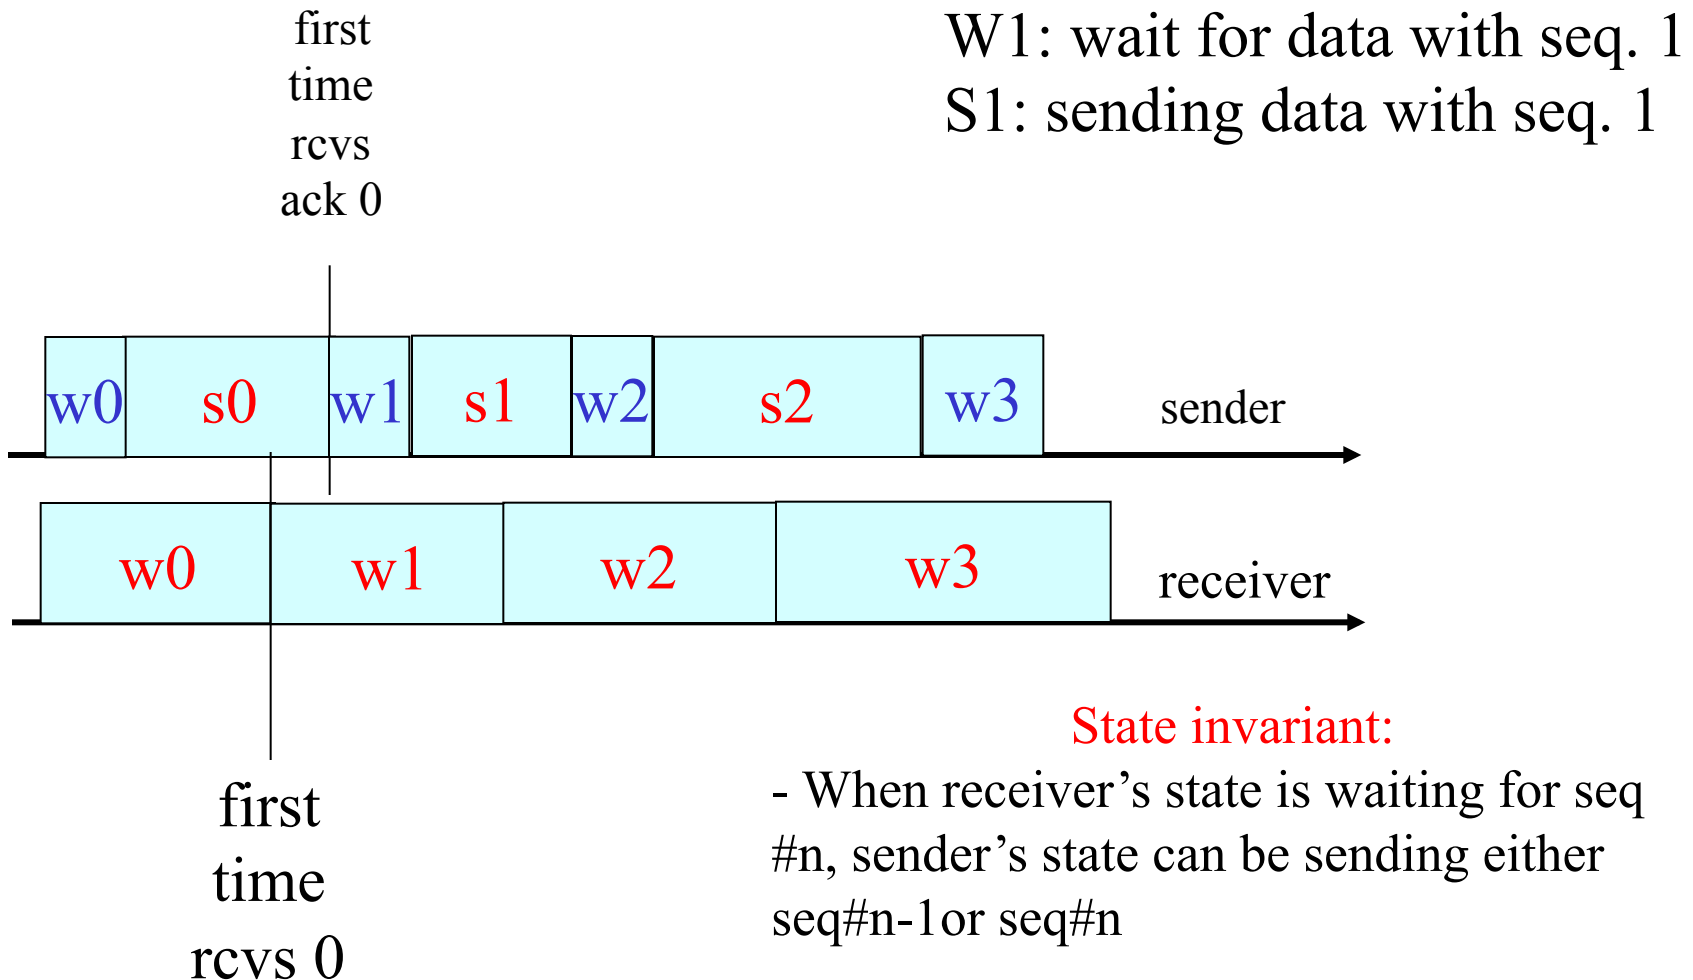


# Another Look at rdt2.1b





# State Relationship of rt2.1b



# rdt2.1b: Summary

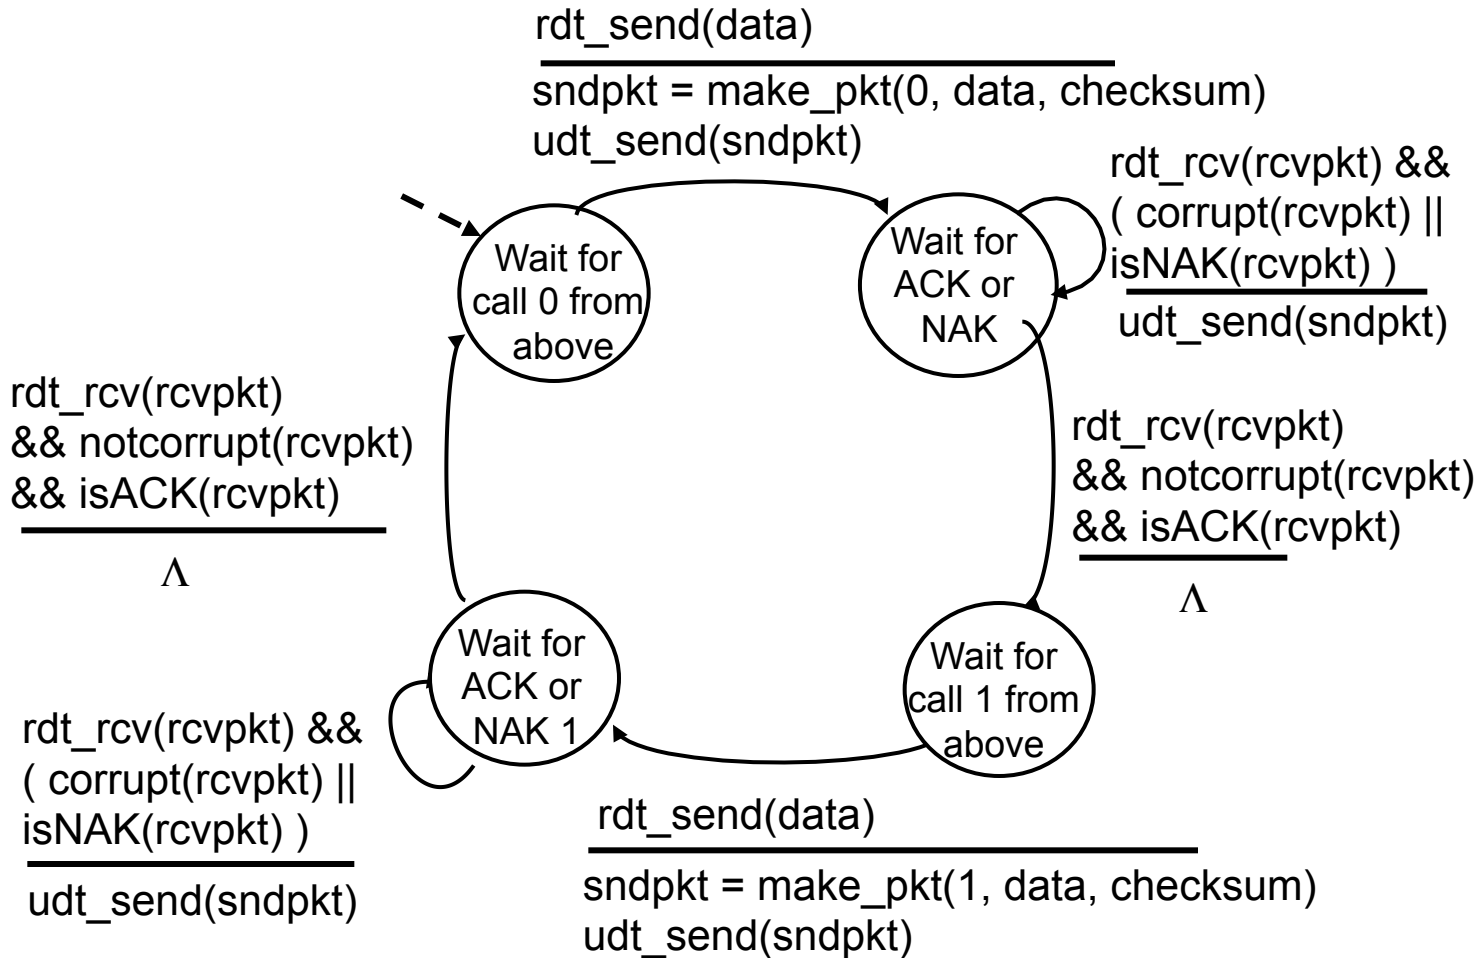
## Sender:

- ❑ Seq # added to pkt
- ❑ Must check if received ACK/NAK corrupted

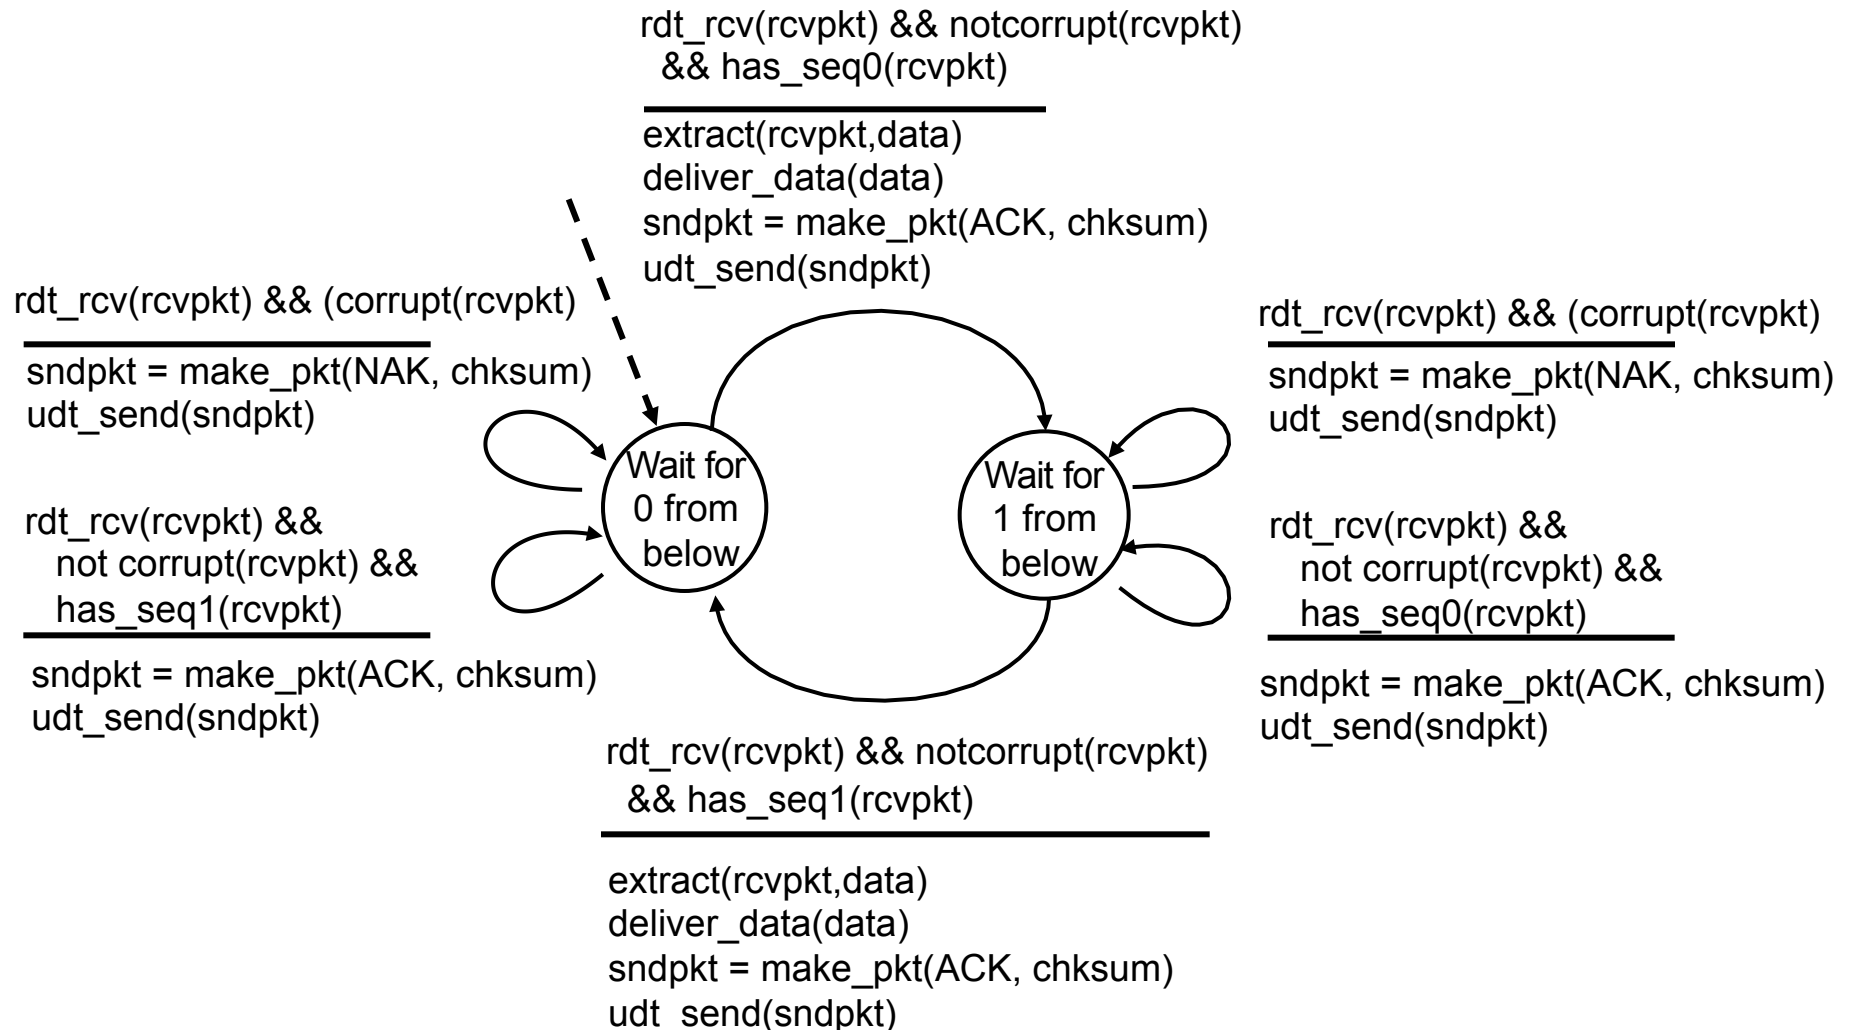
## Receiver:

- ❑ Must check if received packet is duplicate
  - By checking if the packet has the expected pkt seq #

# rdt2.1c: Sender, Handles Garbled ACK/ NAKs: Using 1 bit



# rdt2.1c: Receiver, Handles Garbled ACK/ NAKs: Using 1 bit



# rdt2.1c: Discussion

## Sender:

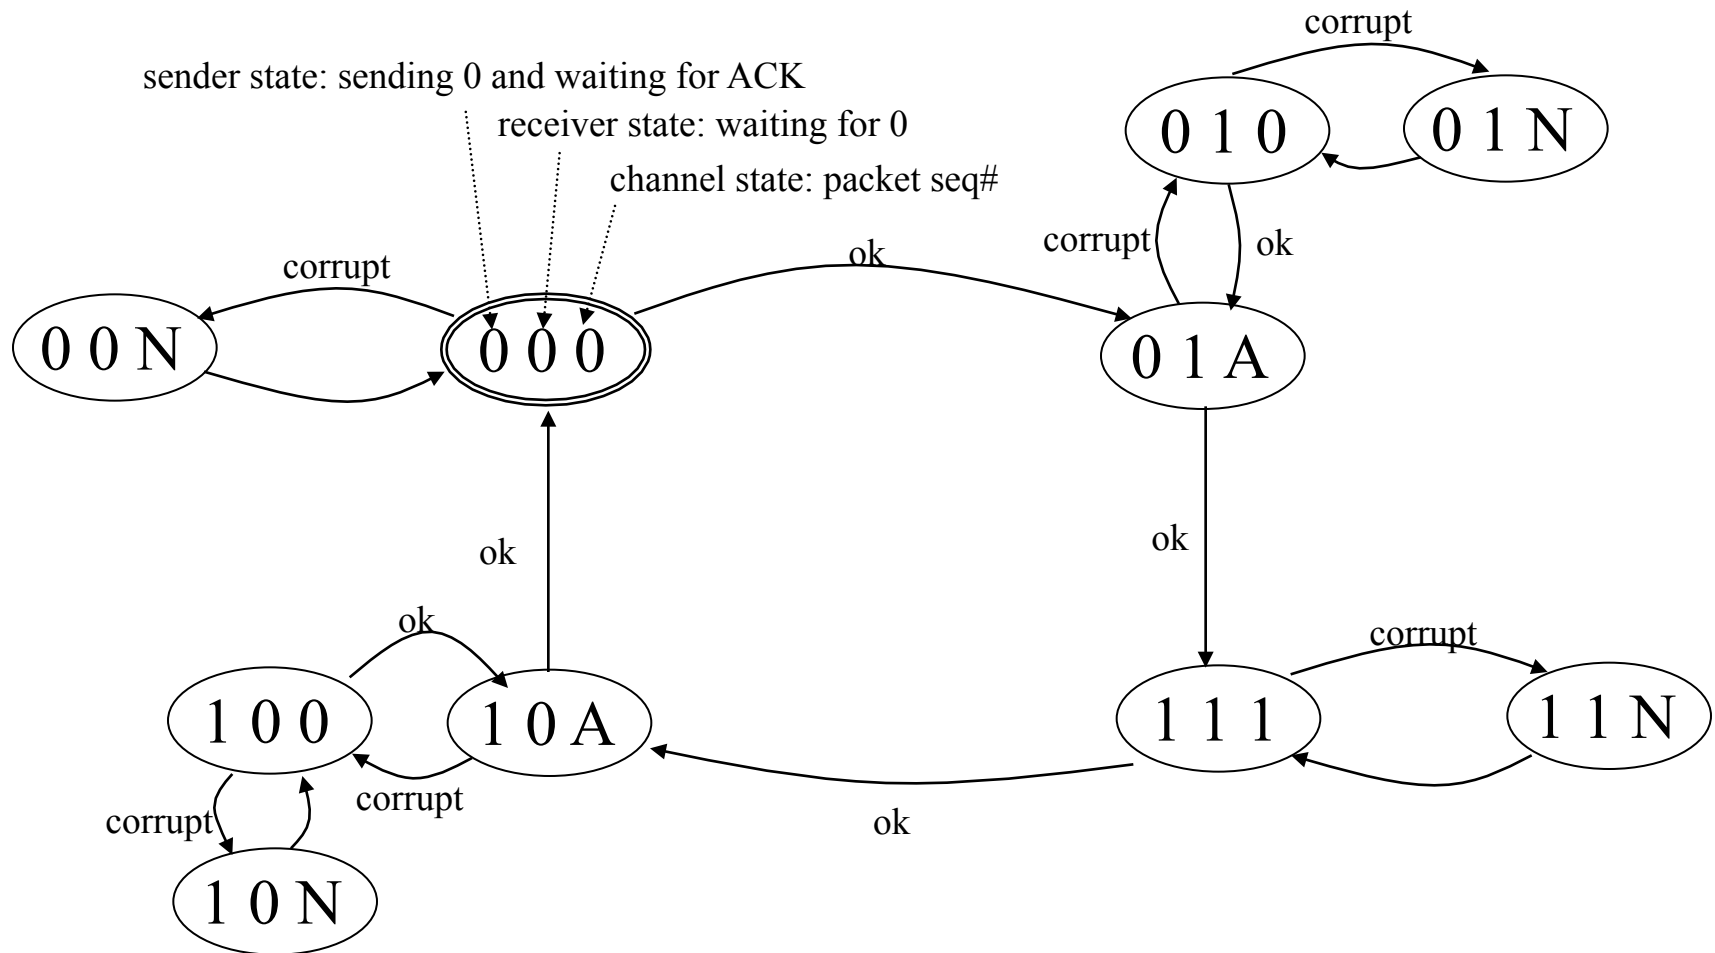
- ❑ State must “remember” whether “current” pkt has 0 or 1 seq. #

## Receiver:

- ❑ Must check if received packet is duplicate
  - State indicates whether 0 or 1 is expected pkt seq #

# A Flavor of Protocol Correctness Analysis: rdt2.1c

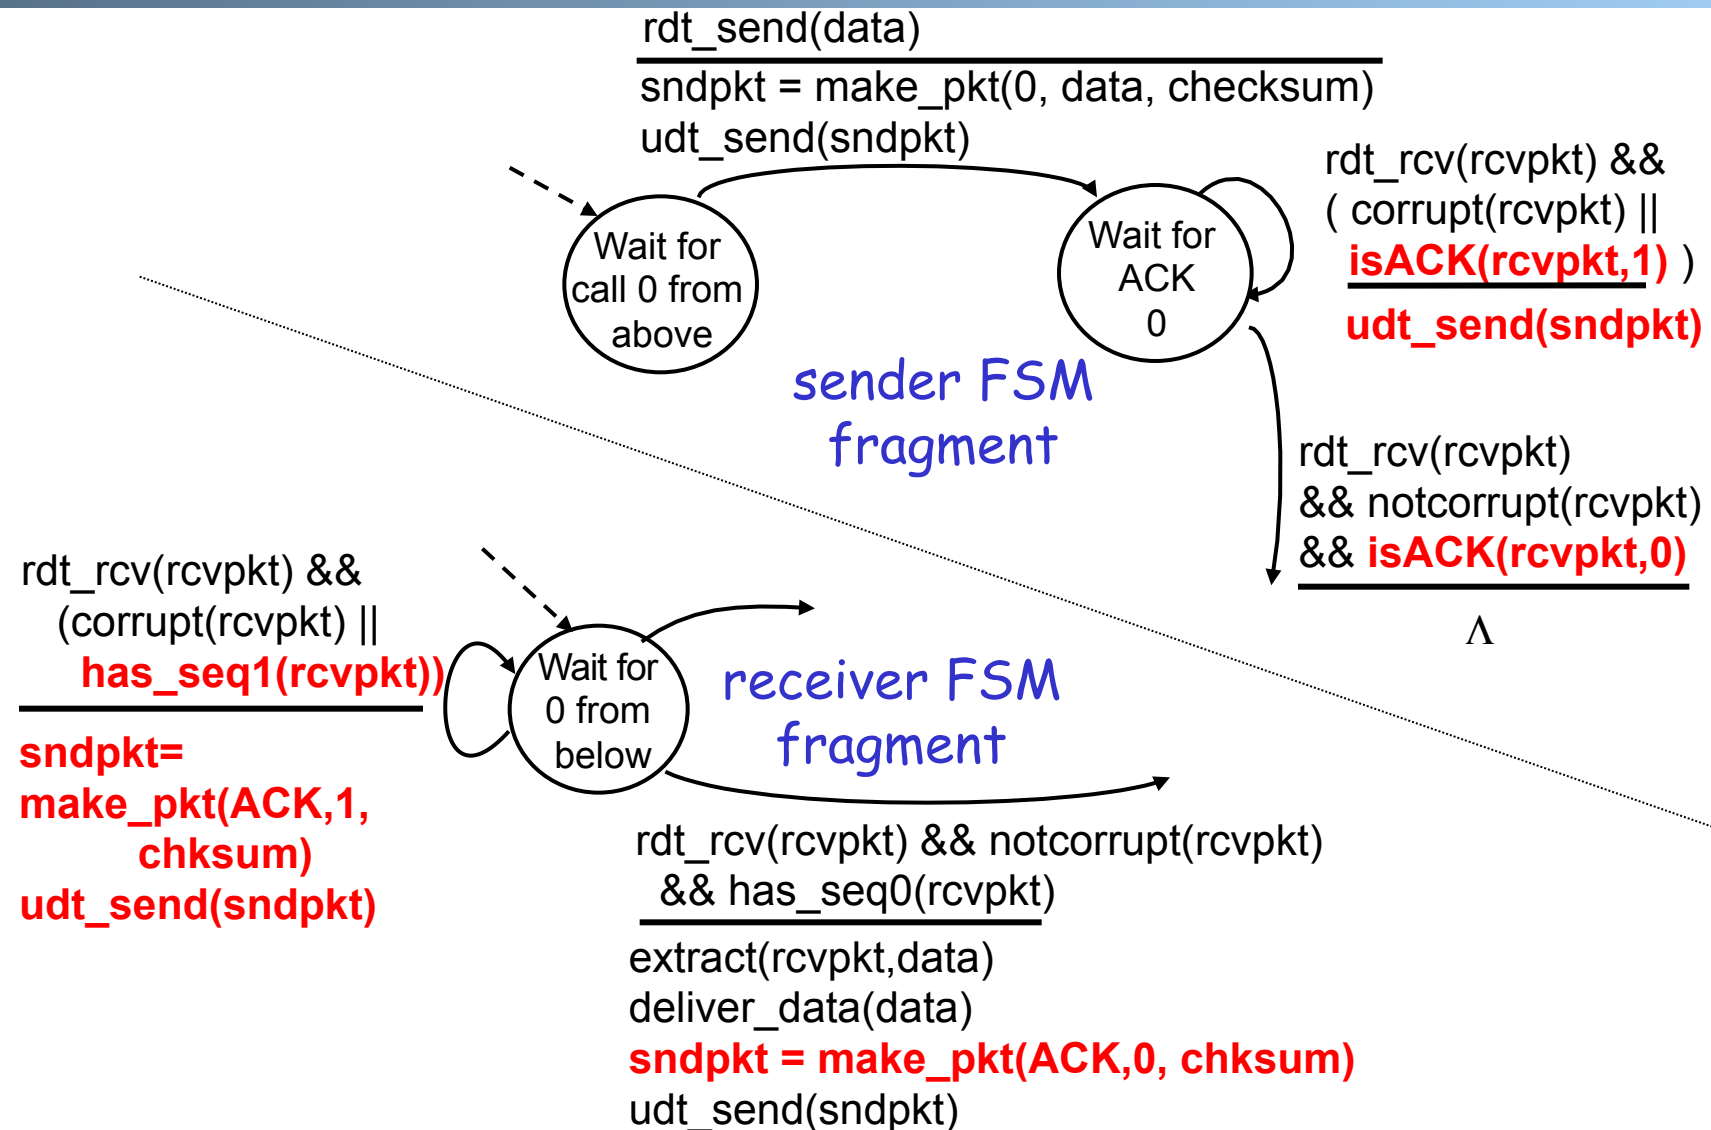
- Combined states of sender and channel
- Assume the sender always has data to send



# rdt2.2: a NAK-free protocol

- ❑ Same functionality as rdt2.1c, using ACKs only
- ❑ Instead of NAK, receiver sends ACK for last pkt received OK
  - Receiver must *explicitly* include seq # of pkt being ACKed
- ❑ Duplicate ACK at sender results in same action as NAK:  
*retransmit current pkt*

# rdt2.2: Sender, Receiver Fragments





# Outline

---

- ❑ Review
- ❑ Overview of transport layer
- ❑ UDP
- Reliable data transfer
  - Perfect channel
  - Channel with bit errors
  - Channel with bit errors and losses

# rdt3.0: Channels with Errors *and* Loss

New assumption: underlying channel can also lose packets (data or ACKs)

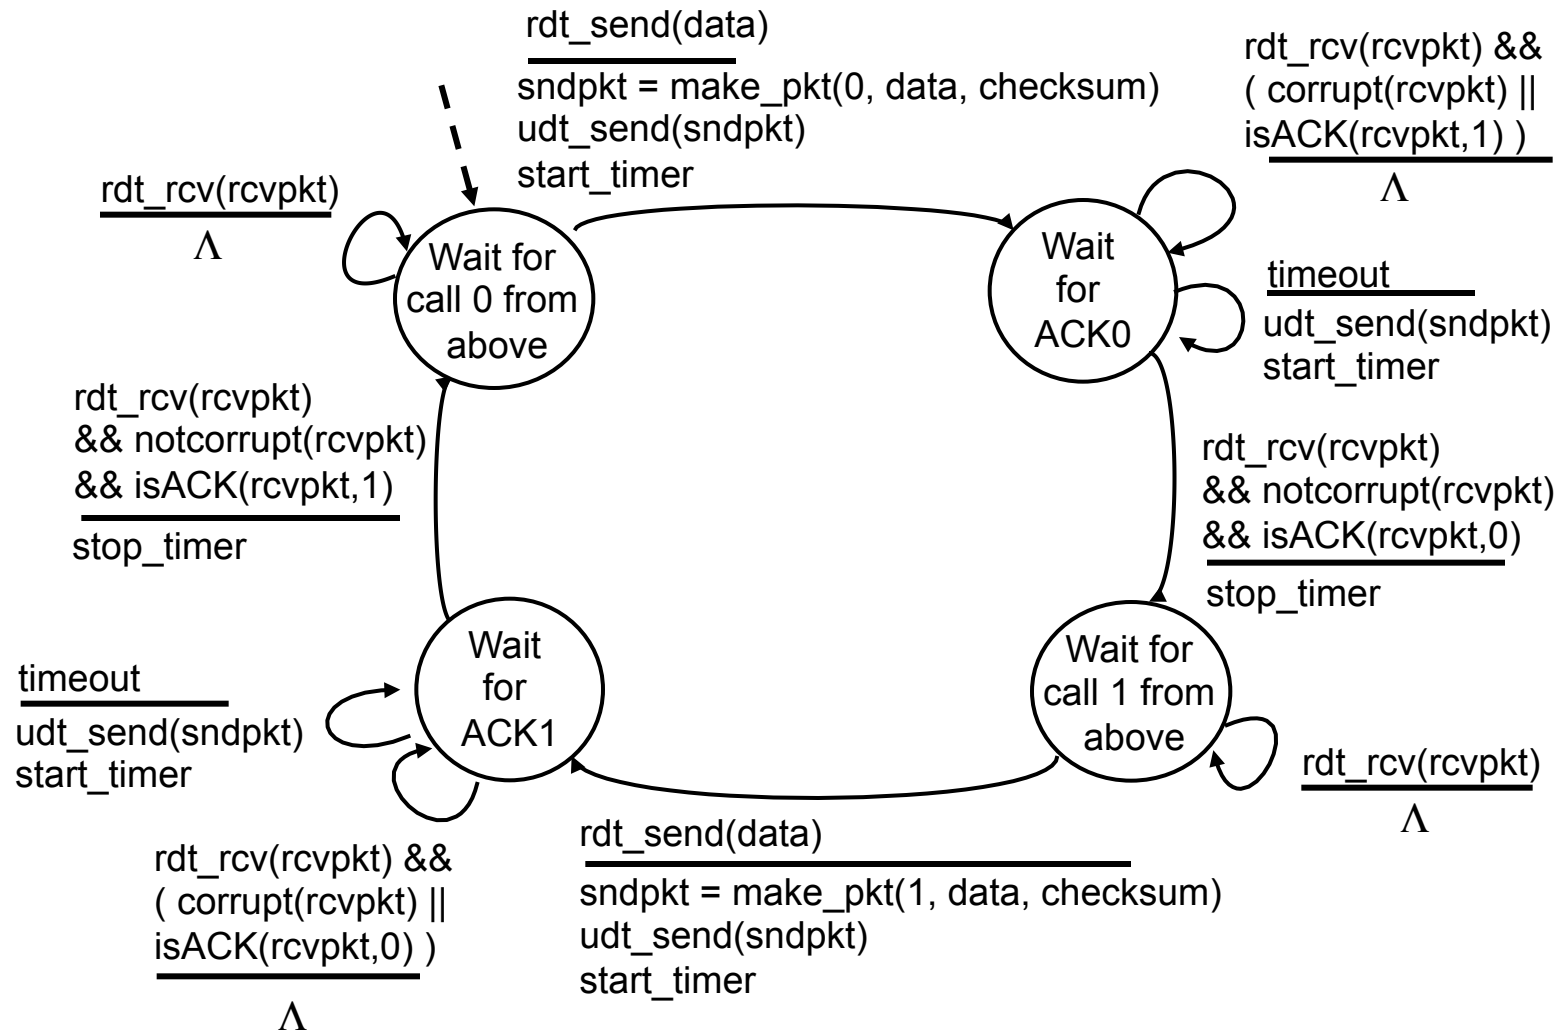
- Checksum, seq. #, ACKs, retransmissions will be of help, but not enough

Q: how to deal with loss?


Approach: sender waits “reasonable” amount of time for ACK

- Requires countdown timer
- Retransmits if no ACK received in this time
- If pkt (or ACK) just delayed (not lost):
  - Retransmission will be duplicate, but use of seq. #'s already handles this
  - Receiver must specify seq # of pkt being ACKed

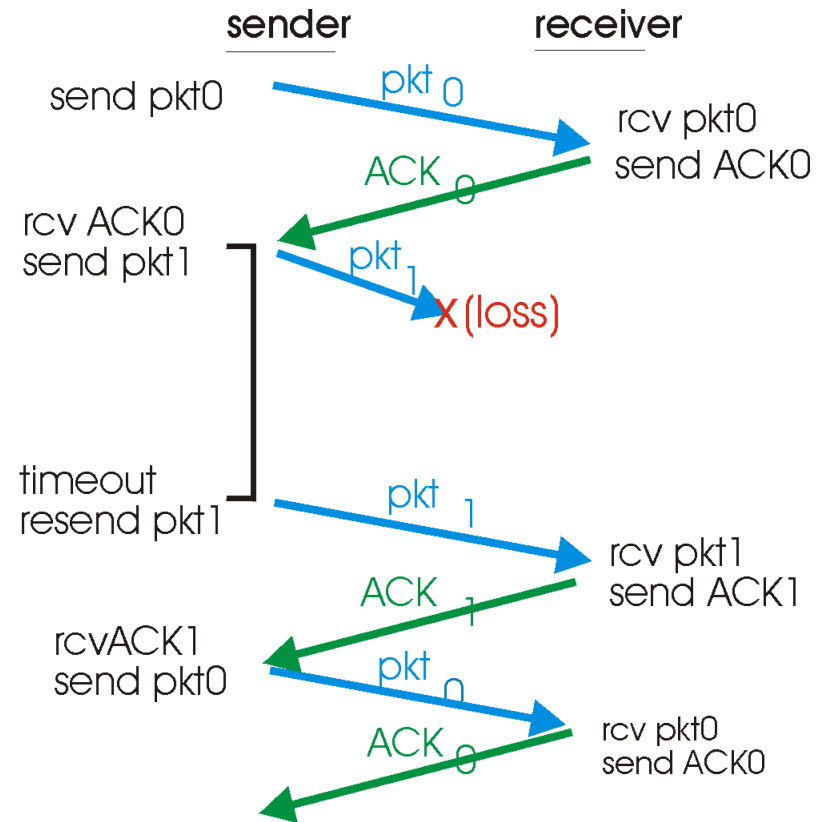
# rdt3.0 Sender



# rdt3.0 in Action



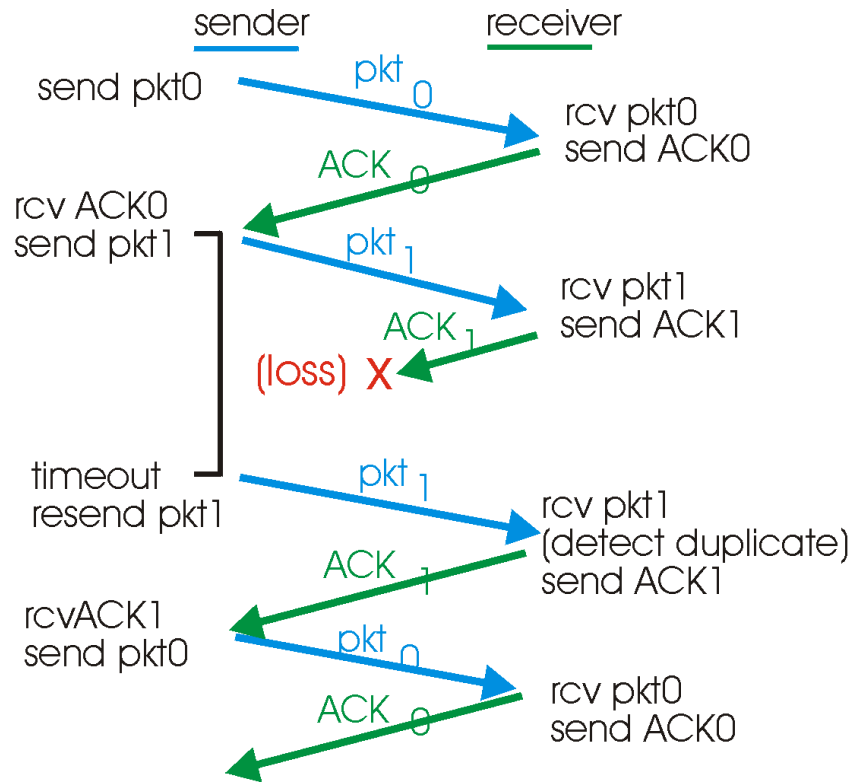
(a) operation with no loss




(b) lost packet

**Question to think about:** How to determine a good timeout value?

# rdt3.0 in Action

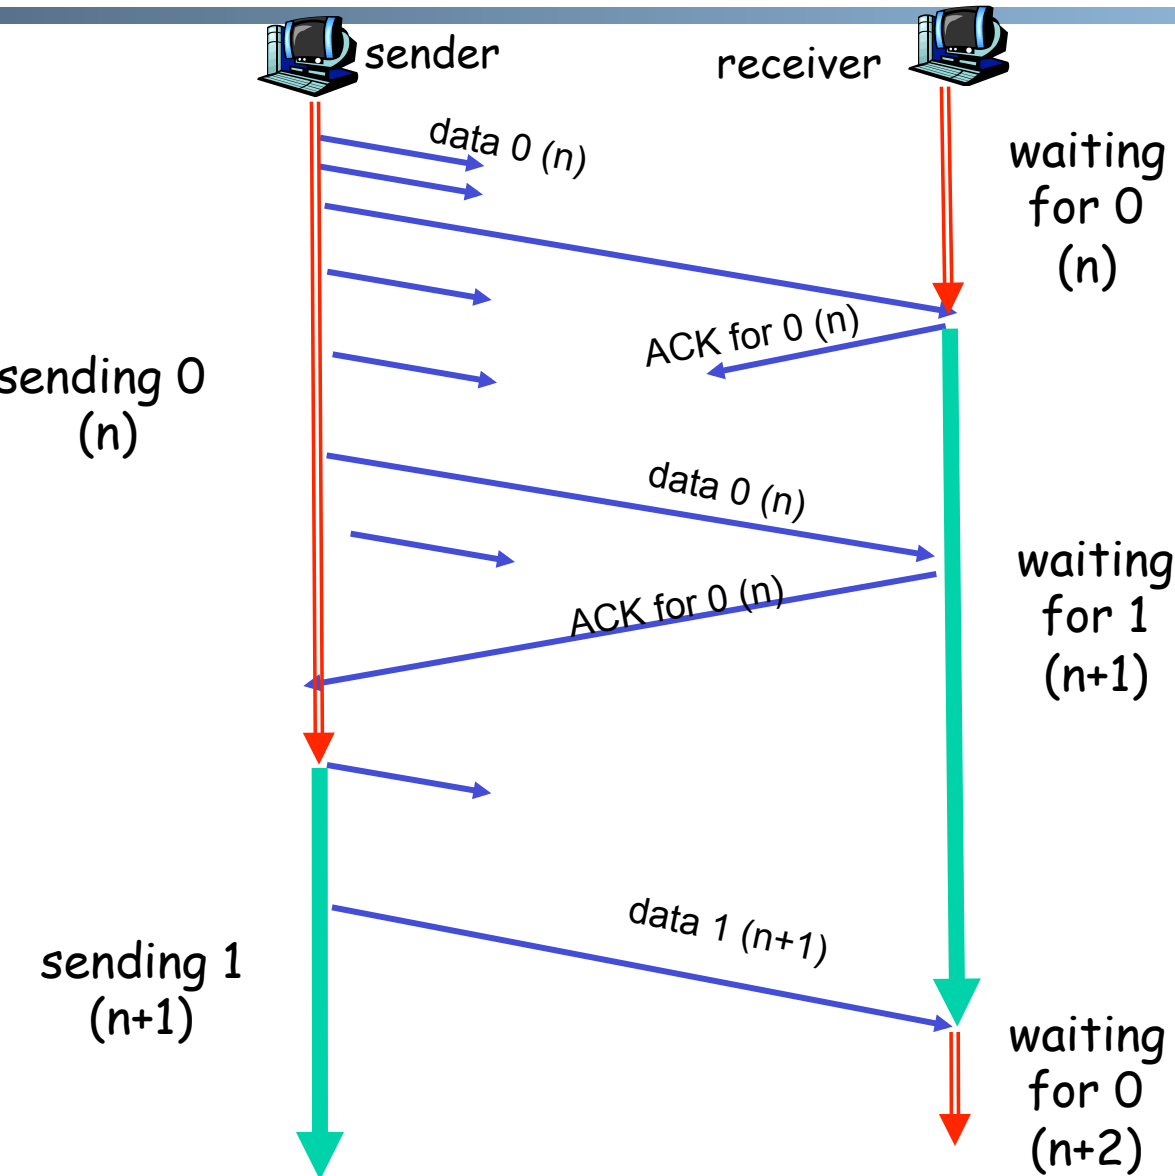


(c) lost ACK



(d) premature timeout

# rdt3.0




## State consistency:

When receiver's state is waiting  $n+1$ , the state of the sender is either sending for  $n+1$  or sending for  $n$

When sender's state is sending for  $n$ , receiver's state is waiting for  $n$  or  $n + 1$

# rdt3.0: Stop-and-Wait Operation



What is  $U_{\text{sender}}$ : **utilization** – fraction of time sender busy sending?

Assume: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet

# Performance of rdt3.0

- ❑ rdt3.0 works, but performance stinks
- ❑ Example: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet:

$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8\text{kb/pkt}}{10^9 \text{ b/sec}} = 8 \text{ microsec}$$

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- 1KB pkt every 30 msec -> 33kB/sec thruput over 1 Gbps link
- Network protocol limits use of physical resources !



# Questions

---

- ❑ How to improve the performance of rdt3.0?
- ❑ How to deal with loss: i.e., how to determine the “right” timeout value?
- ❑ What if there are reordering and duplication?