
Transport Congestion Control: TCP/Reno, Analysis

Quiz

- Explain why TCP's AIMD increases fairness among network flows while maintaining efficiency.

TCP Congestion Control

- ❑ Closed-loop, end-to-end, window-based congestion control
- ❑ Designed by Van Jacobson in late 1980s, based on the AIMD alg. of Dah-Ming Chu and Raj Jain
- ❑ Works well so far: the bandwidth of the Internet has increased by more than 200,000 times
- ❑ Many versions
 - TCP/Tahoe: this is a less optimized version
 - TCP/Reno: many OSs today implement Reno type congestion control
 - TCP/Vegas: not currently used

For more details: see TCP/IP illustrated; or read

http://lxr.linux.no/source/net/ipv4/tcp_input.c for linux implementation

TCP/Reno Congestion Detection

□ Detect congestion in two cases and react differently:

- 3 dup ACKs
- timeout event

Philosophy:


- 3 dup ACKs indicates network capable of delivering some segments
- timeout is “more alarming”

Basic Structure

- ❑ Two “phases”
 - Slow-start: MI
 - Congestion avoidance: AIMD

- ❑ Important variables:
 - *cwnd*: congestion window size
 - *ssthresh*: threshold between the slow-start phase and the congestion avoidance phase

Visualization of the Two Phases



Slow Start: MI

- ❑ What is the goal?
 - Getting to equilibrium gradually but **quickly**
- ❑ Implements the MI algorithm
 - **Double** *cwnd* every RTT until **network congested** → get a rough estimate of the optimal of *cwnd*

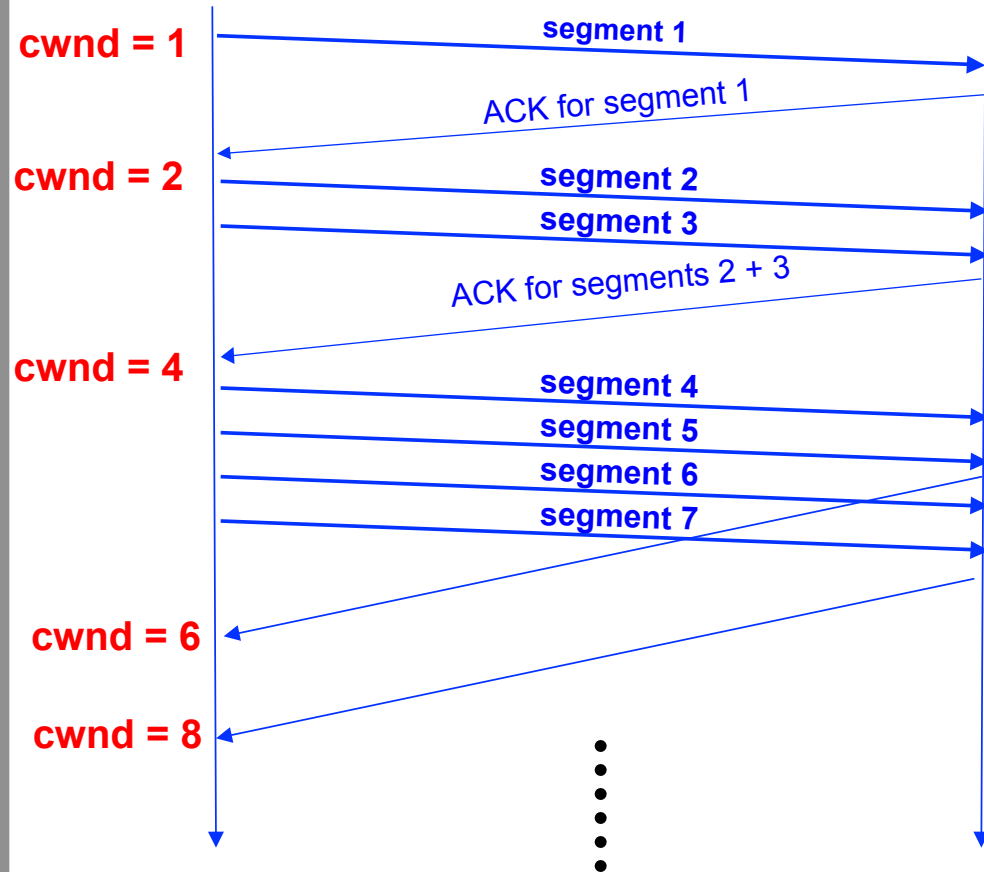
Slow-start

Initially:


`cwnd = 1;`
`ssthresh = infinite (e.g., 64K);`

For each newly ACKed segment:

`if (cwnd < ssthresh)`
 `/* slow start*/`
 `cwnd = cwnd+1;`



Startup Behavior **with** Slow-start



TCP/Reno Congestion Avoidance

- ❑ Maintains equilibrium and reacts around equilibrium
- ❑ Implements the AIMD algorithm
 - Increases window by 1 per round-trip time (how?)
 - Cuts window size
 - To half when detecting congestion by 3 DUP
 - To 1 if timeout
 - If already timeout, *doubles* timeout

TCP/Reno Congestion Avoidance

Initially:

`cwnd = 1;`

`ssthresh = infinite (e.g., 64K);`

For each newly ACKed segment:

`if (cwnd < ssthresh)`

`/* slow start*/`

`cwnd = cwnd + 1;`

`else`

`/* congestion avoidance; cwnd increases (approx.)
by 1 per RTT */`

`cwnd += 1/cwnd;`

Triple-duplicate ACKs:

`/* multiplicative decrease */`

`cwnd = ssthresh = cwnd/2;`


Timeout:

`ssthresh = cwnd/2;`

`cwnd = 1;`

(if already timed out, double timeout value; this is called *exponential backoff*)


TCP/Reno: Big Picture



TD: Triple duplicate acknowledgements

TO: Timeout


A Session



Question: when cwnd is cut to half, why sending rate is not?

TCP/Reno Queueing Dynamics

- Consider congestion avoidance only



There is a filling and draining of buffer process for each TCP flow.

Outline

- ❑ Recap
- ❑ Linear congestion control law
- ❑ TCP/Reno
- *TCP/Reno throughput analysis*

Objective

- ❑ To understand the throughput of TCP/Reno as a function of RTT (RTT), loss rate (p) and packet size
- ❑ We will derive the formula twice, using two setups using two different approaches

TCP/Reno Throughput Modeling

- Given mean packet loss rate p , mean round-trip time RTT , packet size S
- Consider only the congestion avoidance mode (long flows such as large files)
- Assume no timeout
- Assume mean window size is W_m segments, each with S bytes sent in one RTT :


$$\text{throughput} \approx \frac{W_m * S}{RTT} \text{ bytes/sec}$$

Outline

- ❑ Recap
- ❑ Linear congestion control law
- ❑ TCP/Reno
- ❑ TCP/Reno throughput analysis
 - Analysis 1: deterministic

TCP/Reno Throughput Modeling: Relating W with Loss Rate p

□ Consider congestion avoidance only



Assume one packet loss (loss event) per cycle

Total packets send per cycle = $(W/2 + W)/2 * W/2 = 3W^2/8$

Thus $p = 1/(3W^2/8) = 8/(3W^2)$

$$W = \frac{\sqrt{8/3}}{\sqrt{p}} = \frac{1.6}{\sqrt{p}} \Rightarrow \text{throughput} = \frac{S * W_m}{RTT} = \frac{S}{RTT} \frac{3}{4} \frac{1.6}{\sqrt{p}} = \frac{1.2S}{RTT \sqrt{p}}$$

TCP Futures

- ❑ Example: 1500 byte segments, 100ms RTT, want 10 Gbps throughput
- ❑ Requires window size $W = 83,333$ in-flight segments
- ❑ Throughput in terms of loss rate:

$$\frac{1.22 \cdot MSS}{RTT \sqrt{Loss}}$$

- ❑ $\rightarrow Loss = 2 \cdot 10^{-10}$ Wow
- ❑ New versions of TCP for high-speed needed!

Outline

- ❑ Recap
- ❑ Linear congestion control law
- ❑ TCP/Reno
- ❑ TCP/Reno throughput analysis
 - Analysis 1: deterministic
 - Analysis 2: random loss

TCP/Reno Throughput Modeling

$$\Delta W = \begin{cases} \frac{1}{W} & \text{if the packet is not lost} \\ -\frac{W}{2} & \text{if packet is lost} \end{cases}$$

$$\text{mean of } \Delta W = (1-p)\frac{1}{W} + p(-\frac{W}{2}) = 0$$


$$\Rightarrow \text{mean of } W = \sqrt{\frac{2(1-p)}{p}} \approx \frac{1.4}{\sqrt{p}}, \text{ when } p \text{ is small}$$

$$\Rightarrow \text{throughput} \approx \frac{1.4S}{RTT\sqrt{p}}, \text{ when } p \text{ is small}$$

Outline


- ❑ Recap
- ❑ Linear congestion control law
- ❑ TCP/Reno
- ❑ TCP/Reno throughput analysis
 - Analysis 1: deterministic
 - Analysis 2: random loss
- TCP Vegas

TCP/Vegas (Brakmo & Peterson 1994)




- ❑ Idea: try to detect congestion by delay before loss
- ❑ Objective: not to overflow the buffer; instead, try to maintain a *constant* number of packets in the bottleneck queue

TCP/Vegas: Key Question



- How to estimate the number of packets queued in the bottleneck queue?



Background: Little's Law (1961)




□ For any system with no or (low) loss.

□ Assume

○ Mean arrival rate X , mean service time T , and mean number of requests in the system W

□ Then relationship between W , X , and T :


$$W = XT$$



Example: SJTU admits 2500 students each year, and mean time a student stays is 4 years, how many students are enrolled?


Little's Law

$$W = XT$$




$$X = \frac{N}{t} \quad T = \frac{Area}{N} \quad W = \frac{Area}{t}$$

Estimating Number of Packets in the Queue



TCP/Vegas CA algorithm



$$T = T_{\text{prop}} + T_{\text{queueing}}$$

□ Applying Little's Law:

$$x_{\text{vegas}} T = x_{\text{vegas}} T_{\text{prop}} + x_{\text{vegas}} T_{\text{queueing}},$$


where $x_{\text{vegas}} = W / T$ is the sending rate

□ Then number of packets in the queue is

$$\begin{aligned} x_{\text{vegas}} T_{\text{queueing}} &= x_{\text{vegas}} T - x_{\text{vegas}} T_{\text{prop}} \\ &= W - W/T T_{\text{prop}} \end{aligned}$$

TCP/Vegas CA algorithm

maintain a
constant number
of packets in the
bottleneck
buffer



```
for every RTT
{
  if  $W - W/RTT \cdot RTT_{min} < \alpha$  then  $W++$ 
  if  $W - W/RTT \cdot RTT_{min} > \alpha$  then  $W--$ 
}
for every loss
   $W := W/2$ 
```

queue
size

TCP/Vegas Dynamics

$$\Delta w_{\text{RTT}} \approx -(w - xRTT_{\min} - \alpha)$$

$$\Delta w_{\text{unit-time}} = -\left(\frac{w}{RTT} - \frac{x}{RTT} RTT_{\min} - \frac{\alpha}{RTT}\right) = \frac{x}{RTT} RTT_{\min} + \frac{\alpha}{RTT} - x$$

$$\Delta x = \frac{\Delta w_{\text{unit-time}}}{RTT} = \frac{x}{RTT^2} (RTT_{\min} + \frac{\alpha}{x} - RTT)$$

TCP/Reno vs. TCP/Vegas

	TCP/Reno	TCP/Vegas
Congestion signal	loss rate p	queueing delay $T_{queueing}$
Dynamics	$\Delta x = \frac{1}{RTT^2} - p \frac{1}{2} x^2$	$\Delta x = \frac{x}{RTT^2} (RTT_{min} + \frac{\alpha}{x} - RTT)$
Equilibrium	$x_{reno} = \frac{\alpha_{reno}}{RTT \sqrt{p}}$	$x_{vegas} = \frac{\alpha_{vegas}}{T_{queueing}}$

Discussion: Why and why not TCP/Vegas?