
Socket Network Programming

Reading:

Advanced Programming in the UNIX Environment, Chapter 11

Connection-oriented vs. Connectionless

❑ Connectionless protocols

- Use UDP as the transport protocol
- Need to deal with unreliability at the application level

❑ Connection-oriented protocols

- Use TCP as the transport protocol (reliable)
- May use more resources (e.g., three-way handshake, half-open connections)

Stateful vs. Stateless Protocols

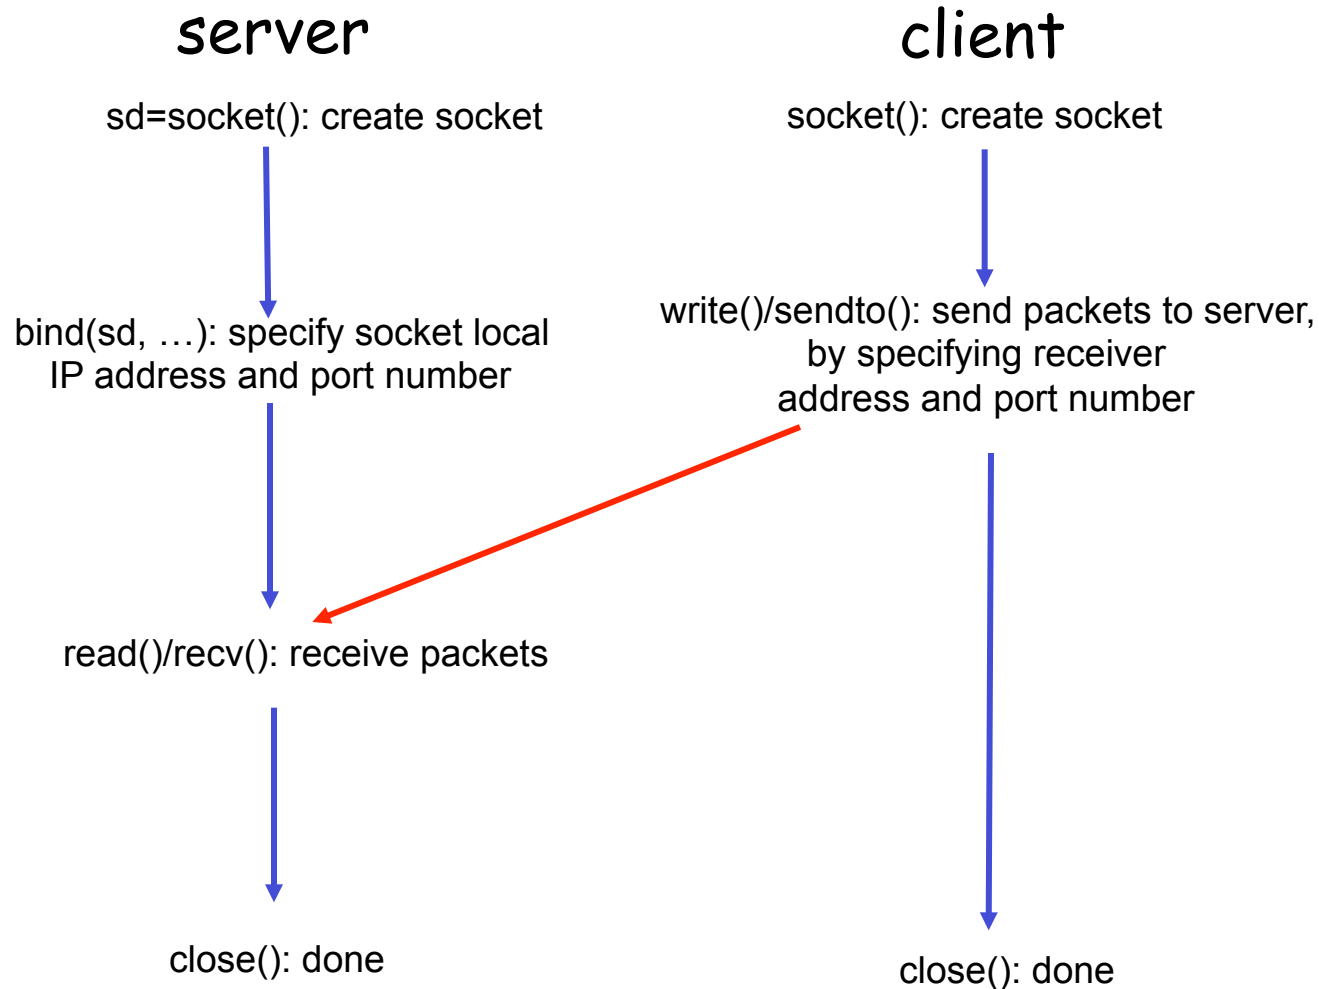
❑ Stateful

- A protocol (and therefore a server) can imply a “history” of previous commands
- The server has to maintain some state associated with the clients
 - Connection-oriented descriptors
 - Client-oriented handles (do not change across connections)
- Need to deal with client/server failures

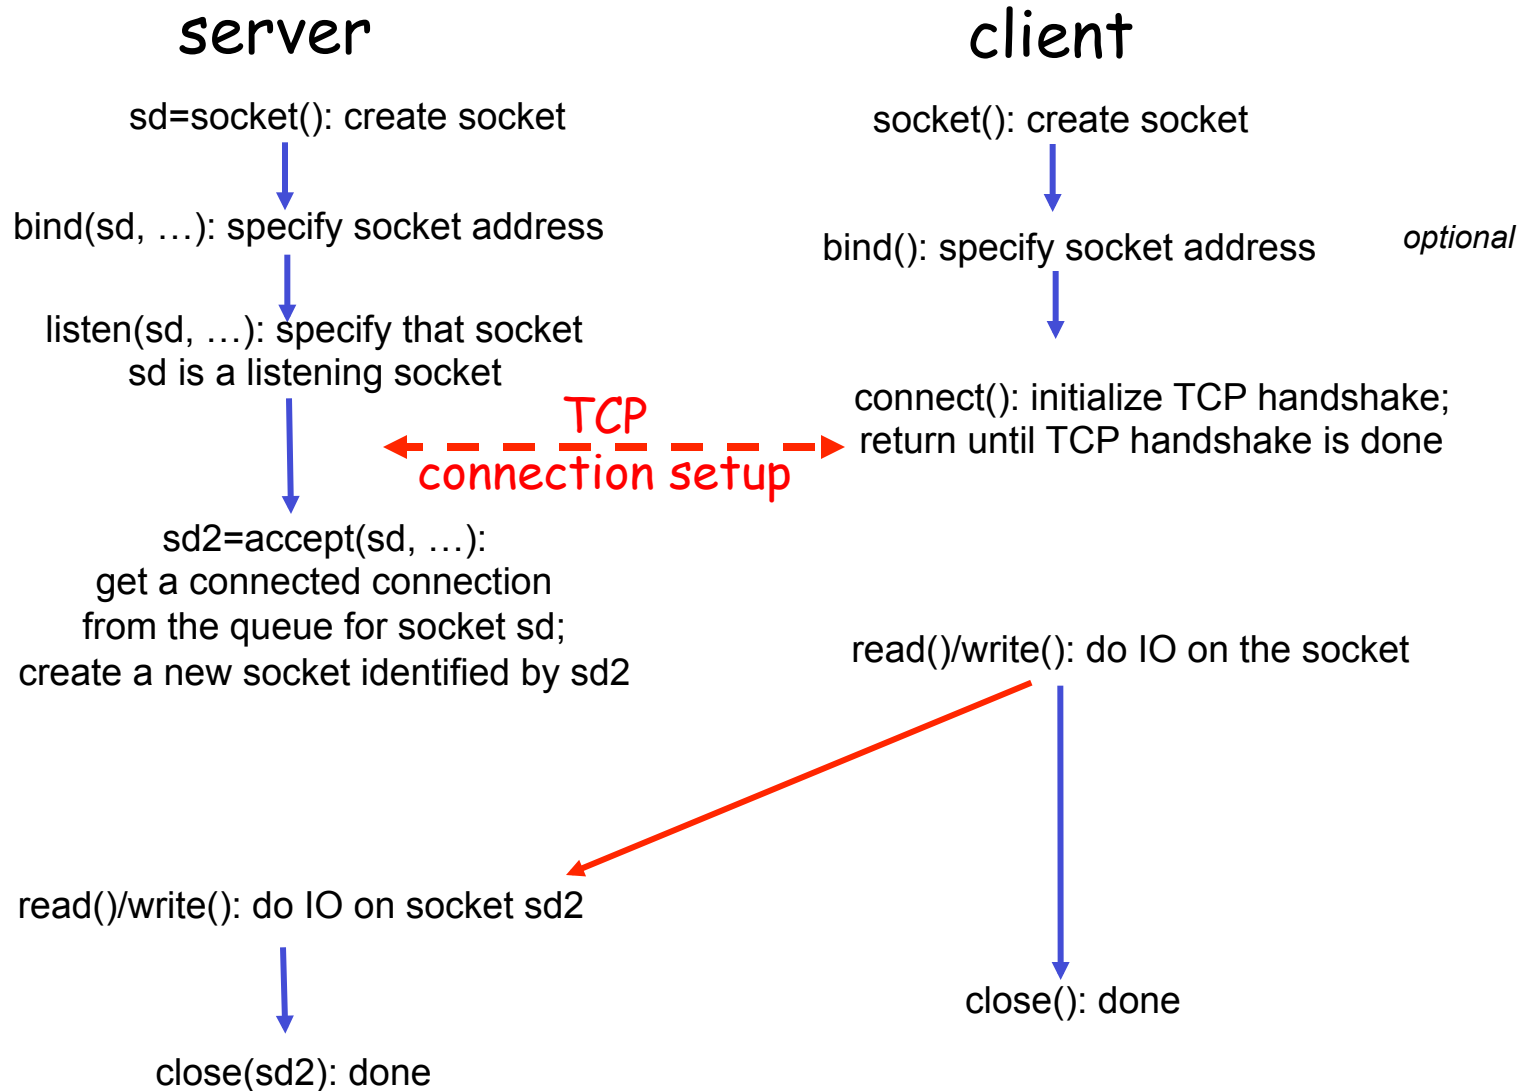
❑ Stateless

- Each protocol interaction is independent
- Might create overhead

Connectionless UDP: Big Picture



Connection-oriented: Big Picture

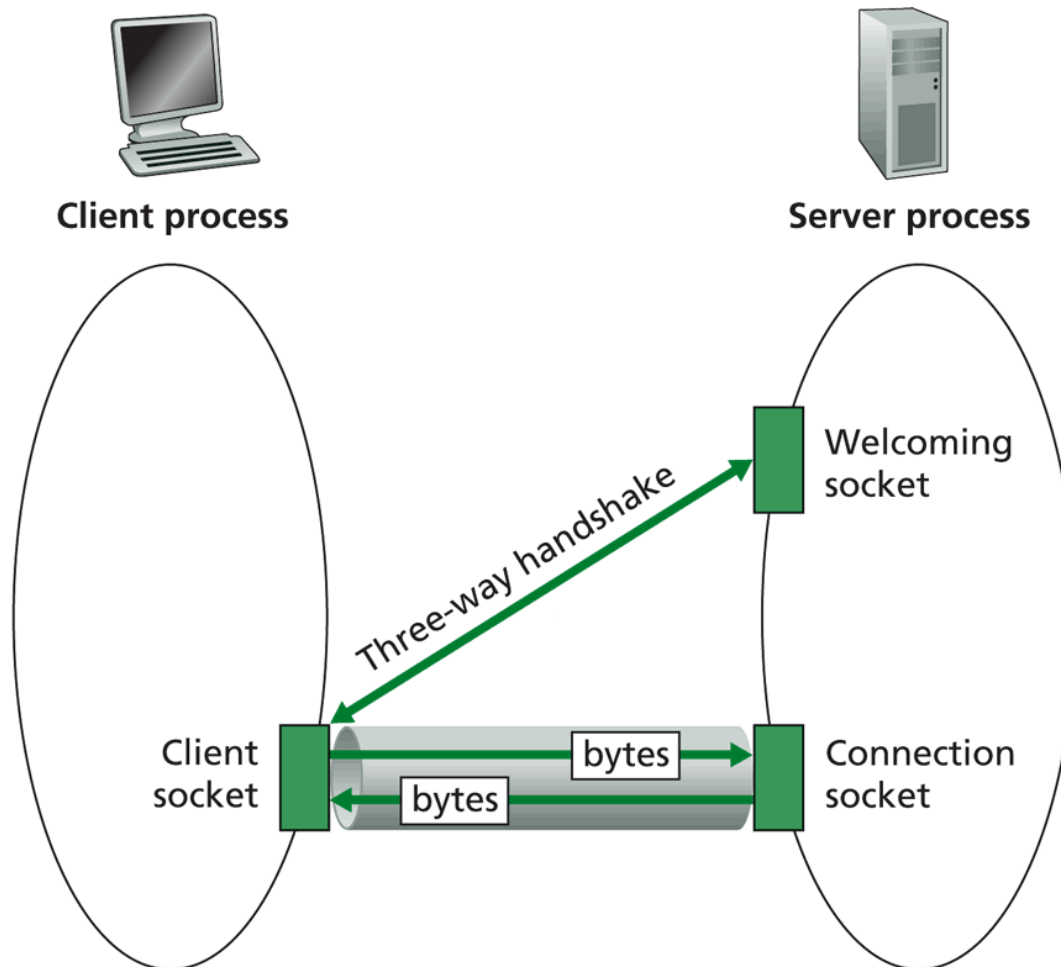


Unix Programming: Mechanism

- ❑ UNIX system calls and library routines (functions called from C/C++ programs)
- ❑ `%man 2 <function name>`
- ❑ A word on style: **check all return codes**

```
if ((code = syscall()) < 0) {  
    perror("syscall");  
}
```

Big Picture: Connection-Oriented TCP



TCP Connection-Oriented Demux

- ❑ TCP socket identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number

- ❑ Host uses all four values to direct segment to appropriate socket
 - Server can easily support many simultaneous TCP sockets: different connections/sessions are automatically separated into different sockets

Socket API

- ❑ API used to access the network
- ❑ A socket is an abstraction of a communication endpoint
 - Can be used to access different network protocols (not only TCP/IP)
 - A socket's characteristics are determined by calling specific functions
 - A socket can be accessed as a file
 - Similar to some I/O APIs, e.g., read, write, close

Create a Socket

- ❑ `int socket(int domain, int type, int protocol)`
 - domain: `PF_UNIX`, `PF_INET`, `PF_INET6`
 - type: `SOCK_STREAM`, `SOCK_DGRAM`
 - protocol: normally 0, can be other values if there are multiple protocols available (`IPPROTO_TCP`, `IPPROTO_UDP`)
 - Return value: the file descriptor of the socket or -1 in case of errors
- ❑ UDP and TCP are accessed using the same primitives but the semantics are different

Creating Sockets Examples

```
if ((sockfd = socket(AF_INET,SOCK_STREAM,0)) {  
    perror("socket");  
    exit(1);  
}
```

Socket Descriptors

- ❑ Similar to file descriptor
 - Internal data structure specifies
 - Protocol used (or family – PF_INET)
 - Type of service (connection-less or connection-oriented)
 - IP addresses involved
 - Ports involved
- ❑ Per-process resource

Bind Sockets to Addresses

- ❑ A socket can be bound to a specific IP address and port
- ❑ `int bind(int sockfd, struct sockaddr *my_addr, socklen_t addrlen)`
 - `sockfd`: socket descriptor
 - `my_addr`: socket address (usually of `sockaddr_in` type)
 - `addrlen`: address structure length
 - Return value: 0 for success, -1 in case of error
- ❑ Same semantics for TCP and UDP

Socket Addresses

```
struct sockaddr {  
    u_short sa_family; /* family */  
    char sa_data[14]; /* address data */  
};
```

```
struct sockaddr_in { /* a TCP endpoint */  
    u_int16_t sin_family; /* address family: AF_INET */  
    u_int16_t sin_port; /* port in network byte order */  
    struct in_addr sin_addr; /* internet address */  
};
```

```
struct in_addr {  
    u_int32_t s_addr; /* address in network byte order */  
};
```

Internet Addresses and Ports

❑ sin_port

- 16 bits
- 0-1024 reserved for system
- well-known ports are important
- If you specify 0, the OS picks a port

❑ s_addr

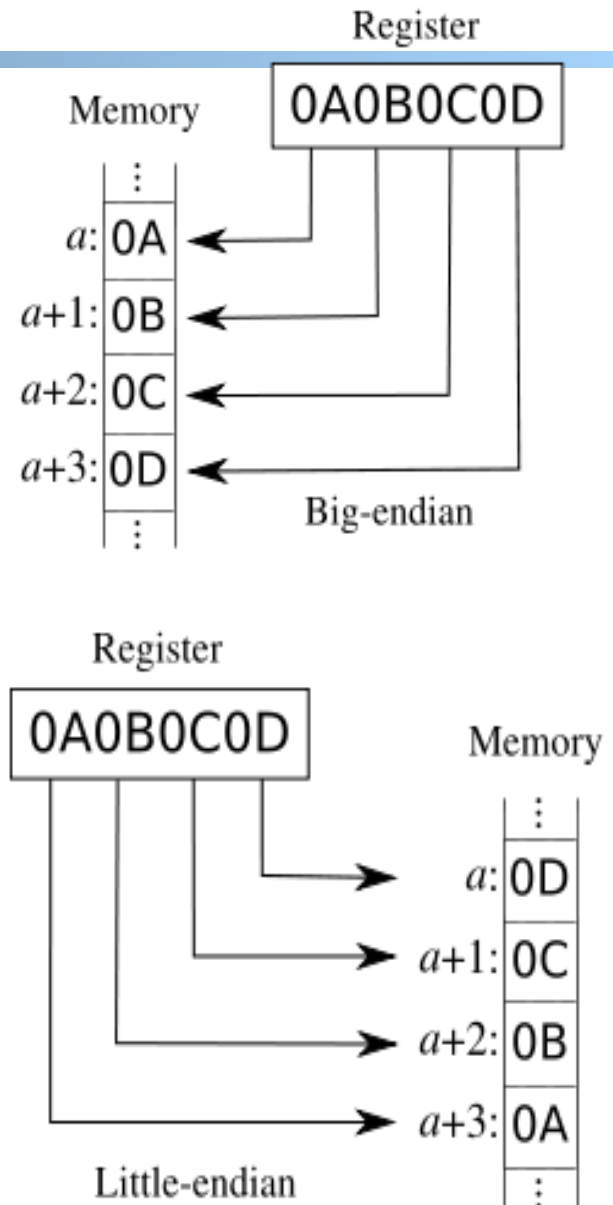
- 32 bits
- INADDR_ANY for any local interface address

Network Byte Order

- ❑ *Big-endian* is network byte order
 - Most-Significant-Byte at lowest memory location
 - E.g., ARM, PowerPC (not the PPC970/G5), DEC Alpha, SPARC V9, MIPS, PA-RISC and IA64

- ❑ *Little-endian*
 - Least-Significant-Byte at lowest memory location
 - E.g., 6502, Z80, x86, and VAX

- ❑ htons, htonl, ntohs, ntohl



Internet Addresses and Ports: Example

```
struct sockaddr_in myaddr;

bzero( (char*)myaddr, sizeof(myaddr) );
myaddr.sin_family = AF_INET;
myaddr.sin_port = htons(80); /* bind to HTTP port*/
myaddr.sin_addr.s_addr = htonl(INADDR_ANY); /* any address*/

if ( (bind(sockfd, (struct sockaddr*)&myaddr, sizeof(myaddr)) < 0 ) {
    perror("bind");
    exit(1);
}
```

Passive Connection-oriented Sockets

- ❑ If a socket is used to receive TCP connections, it is necessary to bind it to an address and to specify that the socket is passive
- ❑ `int listen(int s, int backlog)`
 - s: socket descriptor
 - backlog: maximum length of the queue of pending connections (used to be the number of open/half-open connections, now only the open ones that are ready to be accepted are counted)
 - Return value: 0 in case of success, -1 in case of error

Accept Connections

❑ `int accept(int s, struct sockaddr *addr, socklen_t *addrlen)`

- `s`: socket descriptor
- `addr`: structure that will be filled with the parameters of the client (maybe NULL)
- `addrlen`: length of the structure in input, it is filled with the actual size of the data returned
- Return value: **a new socket file descriptor** in case of success, -1 in case of errors

Accept TCP Connections

- ❑ A call to `accept()` blocks the caller until a request is sent
- ❑ When a connection is made, `accept()` returns a new socket, associated with a specific client (virtual circuit)
- ❑ A virtual circuit is identified by:
 - `<srcIP, srcPort, dstIP, dstPort>`
- ❑ There are no two identical virtual circuits at one time in the whole Internet

Create Connections

- ❑ Clients use `connect()` to open a connection to a specific IP address/port
- ❑ `int connect(int sockfd, const struct sockaddr *serv_addr, socklen_t addrlen)`
 - `sockfd`: socket descriptor
 - `serv_addr`: destination address
 - `addrlen`: size of the structure
 - Return value: 0 in case of success, -1 in case of errors

Connecting with TCP/UDP

❑ TCP

- `connect()` starts the three-way handshake

❑ UDP

- Nothing really happens, but the socket can only be used to send/receive datagrams to/from the specified address

Read/Write to a Socket

- ❑ `read()/write()` of the file interface for connected-oriented
- ❑ Socket specific system call
 - `send()/sendto()/sendmsg()`
 - `recv()/recvfrom()/recvmsg()`

Read from a socket by using read()

```
#include <unistd.h>
```

```
ssize_t read(int sockfd, void *buf, size_t count);
```

- read **up to** count from the socket
- return value: -1 if an error occurs; 0 if end of file; otherwise number of bytes read
- what are the possible outcomes of this system call?

Write to a socket by using write()

```
#include <unistd.h>
```

```
ssize_t write(int sockfd, const void *buf, size_t  
count);
```

- write **up to** count to the socket
- return value: -1 if an error occurs; otherwise number of bytes write
- what are the possible outcomes of this system call?

Send to a Socket

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int send(int sd, const void *msg, size_t len, int flags);
```

```
int sendto(int sd, const void *msg, size_t len, int flags, const  
    struct sockaddr *to, socklen_t tolen);
```

```
int sendmsg(int sd, const struct msghdr *msg, int flags)
```

- return value: -1 if an error occurs; otherwise the number of bytes sent

sendmsg(): scatter and collect

```
struct msghdr {  
    void          *msg_name;    // peer address  
    socklen_t     msg_namelen;  // address length  
    struct iovec  *msg_iov;      // io vector  
    size_t        msg_iovlen;   // io vector length  
    void          *msg_control;  
    socklen_t     msg_controllen;  
    int           msg_flags;  
};  
  
struct iovec {  
    void          *iov_base;  
    size_t        iov_len;  
};
```

Receive from a Socket

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int recv(int sd, void *buf, size_t len, int flags);
```

```
int recvfrom(int sd, void *buf, size_t len, int flags, struct sockaddr  
    *from, socklen_t fromlen);
```

```
int recvmsg(int sd, struct msghdr *msg, int flags);
```

- return value: -1 if an error occurs; otherwise the number of bytes received
- will block the process if there is no data

Close Sockets

❑ `int close(int fd)`

- `fd`: socket descriptor
- Return value: 0 in case of success, -1 in case of errors

❑ `int shutdown(int s, int how)`

- Can be used to close a socket partially
- `s`: connected socket
- `how`:
 - `SHUT_RD`, further receptions are disallowed
 - `SHUT_WR`, further transmissions are disallowed
 - `SHUT_RDWR`, further receptions and transmissions are disallowed
- Return value: 0 in case of success, -1 in case of errors

Closing Semantics

❑ TCP

- Close: FIN/ACK in both directions
- Shutdown: FIN/ACK in one direction

❑ UDP

- Close: Don't send anything. Just deallocate structure

Data Exchange in TCP/UDP

□ TCP

- recv/read will return a chunk of data
- Not necessarily the data sent by means of a single send/write operation on the other side

□ UDP

- recv/read will always return a datagram
- If message size > buffer, fills buffer, **discards rest**

Support Routines: Address from and to String Formats

```
#include <sys/types.h>
#include <sys/socket.h>
#include <arpa/inet.h>
```



presentation format
to network format

```
int inet_pton(int af, const char *src, void *dst);
```

return value: **positive** if successful

```
const char *inet_ntop(int af, const void *src, char *dst, size_t  
cnt);
```

return value: NULL is error

Note: inet_addr() deprecated!

Support Routines: Network/Host Order

```
#include <netinet/in.h>
```

```
unsigned long int htonl(unsigned long int hostlong);
```

```
unsigned short int htons(unsigned short int hostshort);
```

```
unsigned long int ntohl(unsigned long int networklong);
```

```
unsigned short int ntohs(unsigned short int networkshort);
```

DNS Service

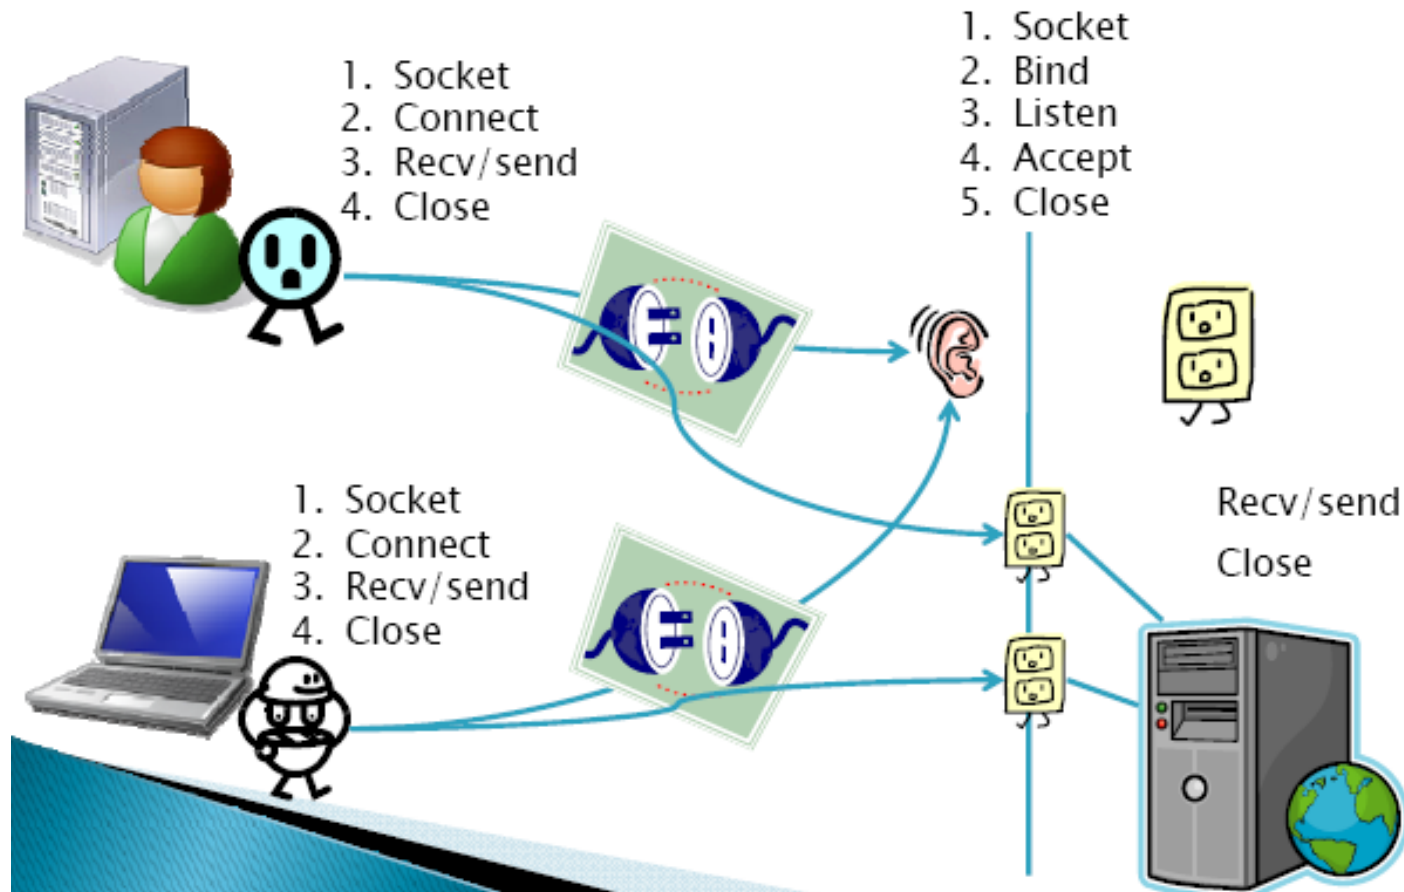
```
#include <netdb.h>
extern int h_errno;

struct hostent *gethostbyname(const char *name);

struct hostent {
    char *h_name;        // official name
    char **h_aliases;    // a list of aliases
    int    h_addrtype;
    int    h_length;
    char **h_addr_list;
}

#define h_addr h_addr_list[0]
- return value: NULL if fails
```

Summary: TCP Client and Server



Summary: Berkeley Sockets

Primitive	Meaning
Socket	Create a new communication end point
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections; given queue size
Accept	Block the caller until a connection attempt arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Receive	Receive some data from the connection
Close	Release the connection

Summary: Socket Programming

❑ \$ man 2 <name>


❑ System calls (functions)

- `int socket(int domain, int type, int protocol);`
- `int bind(int sd, struct sockaddr *my_addr, socklen_t addrlen);`
- `int listen(int sd, int backlog);`
- `int connect(int sd, const struct sockaddr *serv_addr, socklen_t addrlen);`
- `int accept(int sd, struct sockaddr *peer_addr, socklen_t addrlen);`
- `read(int sockfd, void *buf, size_t count);`
- `write(int sockfd, const void *buf, size_t count);`

Asynchronous Network Programming

(C/C++)

A Relay TCP Client: telnet-like Program



Method 1: Process and Thread

❑ Process

- `fork()`
- `waitpid()`

❑ Thread: light weight process

- `pthread_create()`
- `pthread_exit()`

pthread

```
void main()
{
    char recvline[MAXLINE + 1];
    ss = new socketstream(sockfd);

    pthread_t tid;
    if (pthread_create(&tid, NULL,
        copy_to, NULL)) {
        err_quit("pthread_creat()");
    }

    while (ss->read_line(recvline,
        MAXLINE) > 0) {
        fprintf(stdout, "%s\n", recvline);
    }
}
```

```
void *copy_to(void *arg) {
    char sendline[MAXLINE];

    if (debug) cout << "Thread
        create()!" << endl;
    while (fgets(sendline,
        sizeof(sendline), stdin))
        ss->writen_socket(sendline,
            strlen(sendline));

    shutdown(sockfd, SHUT_WR);
    if (debug) cout << "Thread done!"
        << endl;

    pthread_exit(0);
}
```

Method 2: Asynchronous I/O (Select)

- ❑ select: deal with blocking system call

```
int select(int n, fd_set *readfds, fd_set *writefds,  
fd_set *exceptfds, struct timeval *timeout)
```

```
FD_CLR(int fd, fd_set *set);
```

```
FD_ZERO(fd_set *set);
```

```
FD_ISSET(int fd, fd_set *set);
```

```
FD_SET(int fd, fd_set *set);
```

Method 3: Signal and Select

- signal: events such as timeout

Examples of Network Programming

- ❑ Library to make life easier
- ❑ Four design examples
 - TCP Client
 - TCP server using select
 - TCP server using process and thread
 - Reliable UDP
- ❑ Warning: It will be hard to listen to me reading through the code. Read the code.

Example 2: A Concurrent TCP Server Using Process or Thread

- ❑ Get a line, and echo it back
- ❑ Use `select()`
- ❑ For how to use process or thread, see later
- ❑ Check the code at:
<http://jingyu.dyndns.org/~jzhou/courses/F11.Networks/examples-c-socket/tcpserver>
- ❑ Are there potential denial of service problems with the code?

Example 3: A Concurrent HTTP TCP Server Using Process/Thread

- ❑ Use process-per-request or thread-per-request

- ❑ Check the code at:

http://jingyu.dyndns.org/~jzhou/courses/F11.Networks/examples-c-socket/simple_httpd

Example 4: Reliable UDP

- ❑ How to implement timeout?
- ❑ Use SIGALARM
- ❑ Use setjmp()
- ❑ Check the code at:

<http://jingyu.dyndns.org/~jzhou/courses/F10.Networks/examples-c-socket/udptimeout>


Optional Slides

Writing High Performance Servers: Major Issues

- ❑ Many socket/IO operations can cause a process to block, e.g.,
 - `accept`: waiting for new connection;
 - `read` a socket waiting for data or close;
 - `write` a socket waiting for buffer space;
 - I/O `read/write` for disk to finish
- ❑ Thus a crucial perspective of network server design is the concurrency design (non-blocking)
 - for high performance
 - to avoid denial of service
- ❑ Concurrency is also important for clients!



Writing High Performance Servers: Using Multi-Threads

- ❑ Using multiple threads
 - So that only the flow processing a particular request is blocked
 - Java: extends Thread or implements Runnable interface
 - C/C++: use pthread package



Example: a Multi-threaded WebServer, which creates a thread for each request

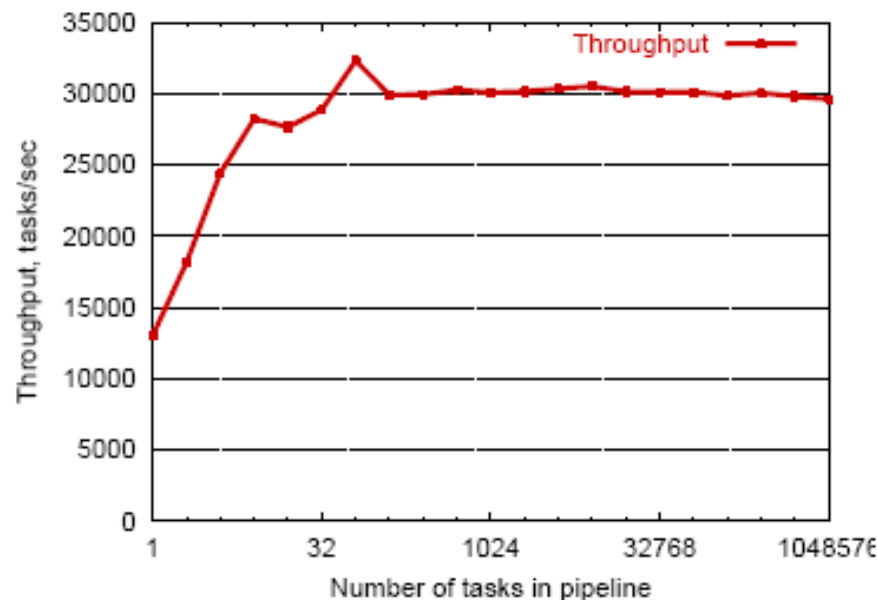

Problems of Multi-Thread Server



(937 MHz x86, Linux 2.2.14, each thread reading 8KB file)

- ❑ High resource usage, context switch overhead, contended locks
- ❑ Too many threads → throughput meltdown, response time explosion
- ❑ In practice: bound total number of threads

Event-Driven Programming



- ❑ Event-driven programming, also called asynchronous i/o
- ❑ Using Finite State Machines (FSM) to monitor the progress of requests
- ❑ Yields efficient and scalable concurrency
- ❑ Many examples: Click router, Flash web server, TP Monitors, etc.
- ❑ Java: asynchronous i/o
 - for an example see: <http://www.cafeaulait.org/books/jnp3/examples/12/>

Problems of Event-Driven Server

- ❑ Difficult to engineer, modularize, and tune
- ❑ Little OS and tool support: ‘ ‘roll your own’ ’
- ❑ No performance/failure isolation between FSMs
- ❑ FSM code can never block (but page faults, garbage collection may still force a block)
 - thus still need multiple threads