

Network Applications: P2P Applications

Summary of Traditional C-S Network Applications


- ❑ How does a client locate a server?
- ❑ Is the application extensible, robust, scalable?



down speed to the clients?
slashdot effect, CNN on 9/11


Objectives of P2P

- ❑ Bypass DNS to access resources!
 - Examples: instant messaging, skype
- ❑ Share the storage *and* bandwidth of individual clients to improve scalability
 - Examples: file sharing and streaming



Peer-to-Peer Computing

- Quickly grown in popularity:
 - Dozens or hundreds of file sharing applications
 - 50-80% Internet traffic is P2P
 - Upset the music industry, drawn college students, web developers, recording artists and universities into court



From ipoque web site; Nov. 2007

What is P2P?

- ❑ But P2P is not new and is probably here to stay
- ❑ Original Internet was a P2P system:
 - The original ARPANET connected UCLA, Stanford Research Institute, UCSB, and Univ. of Utah
 - No DNS or routing infrastructure, just connected by phone lines
 - Computers also served as routers
- ❑ P2P is simply an iteration of scalable distributed systems

P2P Systems


- ❑ File Sharing: BitTorrent, LimeWire, eMule, eDonkey
- ❑ Streaming: PPLive, PPStream, ...
- ❑ Research systems
 - Distributed Hash Tables
 - Content distribution networks
- ❑ Collaborative computing:
 - SETI@Home project (screen saver)
 - Human genome mapping
 - Intel NetBatch: 10,000 computers in 25 worldwide sites for simulations, saved about 500 million

Outline

- P2P

- the lookup problem

The Lookup Problem



find where a particular file is stored
pay particular attention to see its equivalence of DNS


Outline

□ P2P

- The lookup problem
 - Napster

Centralized Database: Napster

- ❑ Program for sharing **music** over the Internet
- ❑ History:
 - **5/99**: Shawn Fanning (freshman, Northeastern U.) founded Napster Online music service, wrote the program in 60 hours
 - **12/99**: first lawsuit
 - **3/00**: 25% UWisc traffic Napster
 - **2000**: est. 60M users
 - **2/01**: US Circuit Court of Appeals: Napster knew users violating copyright laws
 - **7/01**: # simultaneous online users: Napster 160K
 - **9/02**: bankruptcy



03/2000

We are referring to the Napster before closure.

Napster: How Does it Work?

Application-level, client-server protocol over TCP

A centralized index system that maps files (songs) to machines that are alive and with files

Steps:


- ☐ Connect to Napster server
- ☐ Upload your list of files (push) to server
- ☐ Give server keywords to search the full list
- ☐ Select “best” of hosts with answers

Napster Architecture


From Computer Desktop Encyclopedia
© 2004 The Computer Language Co. Inc.

THE ORIGINAL NAPSTER

Napster provided a central directory of users who had files to share.




Napster: Publish



Napster: Search

123.2.0.18



```
search(A)
-->
123.2.0.18
124.1.0.1
```

Query


Reply

124.1.0.1

Where is file A?





Napster: Ping




Napster: Fetch

123.2.0.18



124.1.0.1



fetch



Napster Messages

General Packet Format

[chunksize] [chunkinfo] [data...]

CHUNKSIZE:

Intel-endian 16-bit integer
size of [data...] in bytes

CHUNKINFO: (hex)

Intel-endian 16-bit integer.

00 - login rejected	5B - whois query
02 - login requested	5C - whois result
03 - login accepted	5D - whois: user is offline!
0D - challenge? (nuprin1715)	69 - list all channels
2D - added to hotlist	6A - channel info
2E - browse error (user isn't online!)	90 - join channel
2F - user offline	91 - leave channel

.....

Centralized Database: Napster

- ❑ Summary of features: a hybrid design
 - **Control**: client-server (aka special DNS) for files
 - **Data**: peer to peer
- ❑ Advantages
 - Simplicity, easy to implement sophisticated search engines on top of the index system
- ❑ Disadvantages
 - Application specific (compared with DNS)
 - Lack of robustness, scalability: central search server single point of bottleneck/failure
 - Easy to sue !

Variation: BitTorrent

- ❑ A global central index server is replaced by one tracker per file (called a swarm)
 - Reduces centralization; but needs other means to locate trackers

- ❑ The bandwidth scalability management technique is more interesting
 - More later

Outline

- *P2P*

- *The lookup problem*

- Napster (central query server; distributed data servers)


- *Gnutella*

Gnutella


- ❑ On March 14th 2000, J. Frankel and T. Pepper from AOL's Nullsoft division (also the developers of the popular Winamp mp3 player) released Gnutella
- ❑ Within hours, AOL pulled the plug on it
- ❑ Quickly reverse-engineered and soon many other clients became available: Bearshare, Morpheus, LimeWire, etc.

Decentralized Flooding: Gnutella

- ❑ On startup, client contacts other servents (**server + client**) in network to form interconnection/peering relationships
 - Servent interconnection used to forward control (queries, hits, etc)
- ❑ How to find a resource record: decentralized flooding
 - Send requests to neighbors
 - Neighbors recursively forward the requests




Decentralized Flooding




Decentralized Flooding

send query to neighbors



- Each node forwards the query to its neighbors other than the one who forwards it the query

Background: Decentralized Flooding



- ❑ Each node should keep track of forwarded queries to avoid loop !
 - Nodes keep state (which will time out---soft state)
 - Carry the state in the query, i.e. carry a list of visited nodes

Decentralized Flooding: Gnutella

- ❑ Basic message header
 - Unique ID, TTL, Hops
- ❑ Message types
 - Ping – probes network for other servants
 - Pong – response to ping, contains IP addr, # of files, etc.
 - Query – search criteria + speed requirement of servant
 - QueryHit – successful response to Query, contains addr + port to transfer from, speed of servant, etc.
 - Ping, Queries are flooded
 - QueryHit, Pong: reverse path of previous message

Advantages and Disadvantages of Gnutella

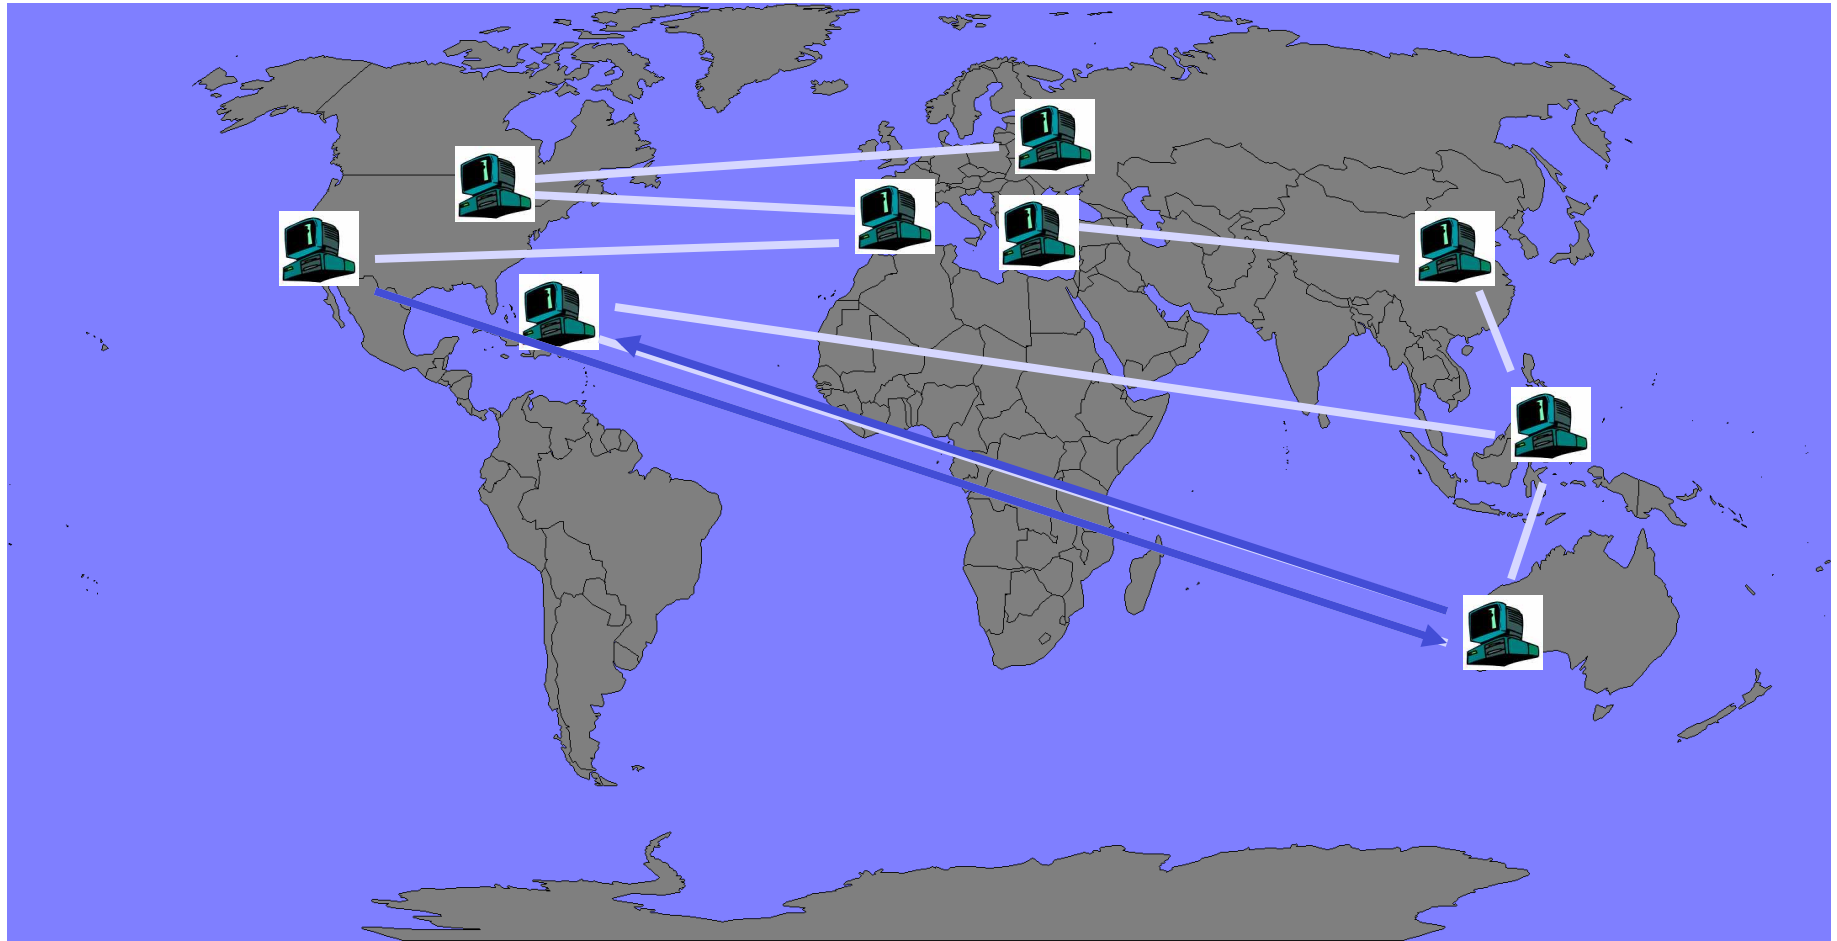
□ Advantages:

- Totally decentralized, highly robust


□ Disadvantages:

- Not scalable; the entire network can be swamped with flood requests
 - Especially hard on slow clients; at some point broadcast traffic on Gnutella exceeded 56 kbps
- To alleviate this problem, each request has a TTL to limit the scope
 - Each query has an initial TTL, and each node forwarding it reduces it by one; if TTL reaches 0, the query is dropped (consequence?)


Aside: Search Time?




Aside: All Peers Equal?




Aside: Network Resilience



Partial Topology



Random 30% die



Targeted 4% die

from Saroiu *et al.*, *MMCN* 2002

Flooding: FastTrack (aka Kazaa)

- ❑ Modifies the Gnutella protocol into two-level hierarchy
- ❑ Supernodes
 - Nodes that have better connection to Internet
 - Act as temporary indexing servers for other nodes
 - Help improve the stability of the network
- ❑ Standard nodes
 - Connect to supernodes and report list of files
- ❑ Search
 - Broadcast (Gnutella-style) search across supernodes
- ❑ Disadvantages
 - Kept a centralized registration → prone to law suits

Outline

➤ *P2P*

➤ *the lookup problem*

- Napster (central query server; distributed data server)
- Gnutella (decentralized, flooding)

➤ *Freenet*

Freenet

❑ History

- Final year project [Ian Clarke](#) , [Edinburgh University](#), Scotland, June, 1999

❑ Goals:

- Totally distributed system without using centralized index or broadcast (flooding)
- Respond adaptively to usage patterns, transparently moving, replicating files as necessary to provide efficient service
- Provide publisher anonymity, security
- Free speech : resistant to attacks – a third party shouldn't be able to deny (e.g., deleting) the access to a particular file (data item, object)

Basic Structure of Freenet

- ❑ Each machine stores a set of files; each file is identified by a unique identifier (called key or id)
- ❑ Each node maintains a “routing table”
 - *id* – file id, key
 - *next_hop node* – where a file corresponding to the id might be available
 - *file* – local copy if one exists

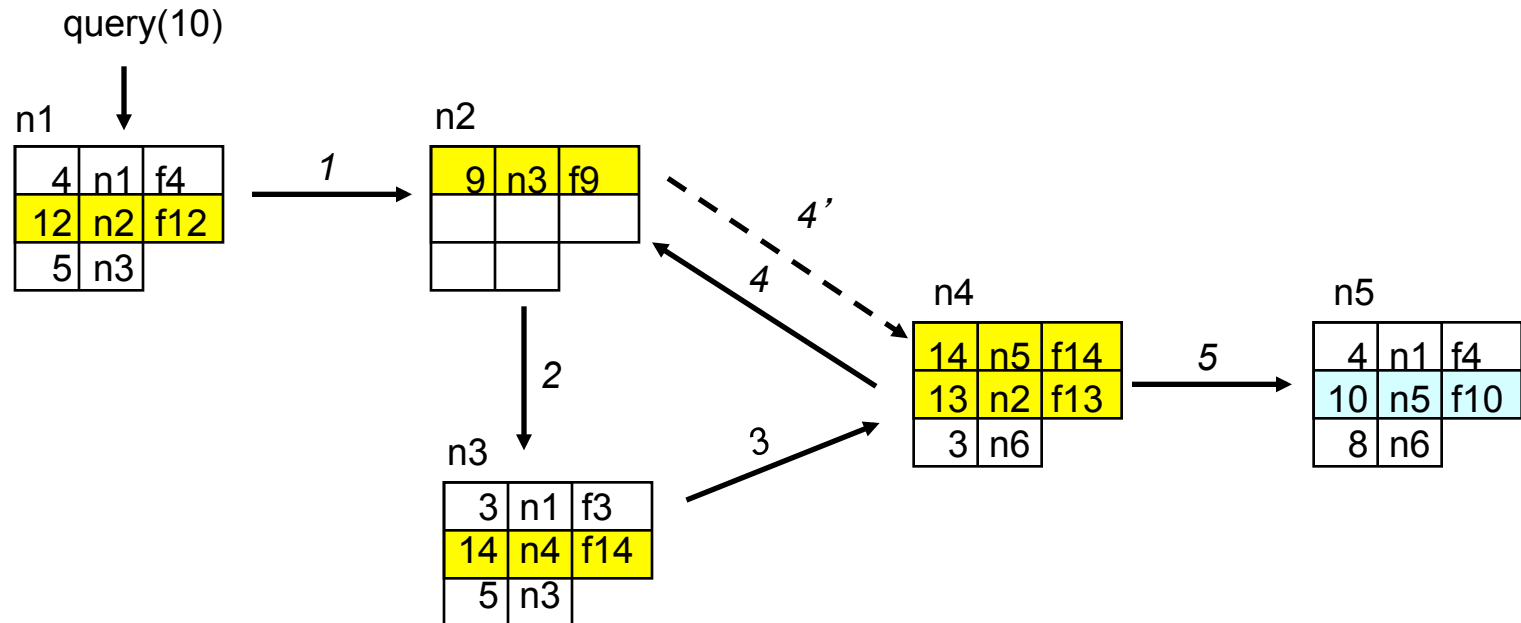
<i>id</i>	<i>next_hop</i>	<i>file</i>
	⋮	
↓	↓	
	⋮	
	⋮	

Query

- ❑ API: $file = query(id)$;
- ❑ Upon receiving a query for file id
 - Check whether the queried file is stored locally
 - If yes, return it
 - If not, forward the query message
 - Key step: search for the “closest” id in the table, and forward the message to the corresponding $next_hop$

<i>id</i>	<i>next_hop</i>	<i>file</i>
	⋮	
↓	↓ ⋮	
	⋮	

Query Example



Beside the routing table, each node also maintains a query table containing the state of all outstanding queries that have traversed it → to backtrack

Query: the Complete Process

- ❑ Search by routing
- ❑ API: *file* = query(*id*);

- ❑ Upon receiving a query for file *id*
 - Check whether the queried file is stored locally
 - If yes, return it; otherwise
 - Check TTL to limit the search scope
 - Each query is associated a TTL that is decremented each time the message is forwarded
 - When TTL=1, the query is forwarded with a probability
 - TTL can be initiated to a random value within some bounds to obscure distance to originator
 - Look for the “**closest**” *id* in the table with an unvisited *next_hop* node
 - If found one, forward the query to the corresponding *next_hop*
 - Otherwise, backtrack
 - Ends up performing a Depth First Search (DFS)-like traversal
 - Search direction ordered by closeness to target

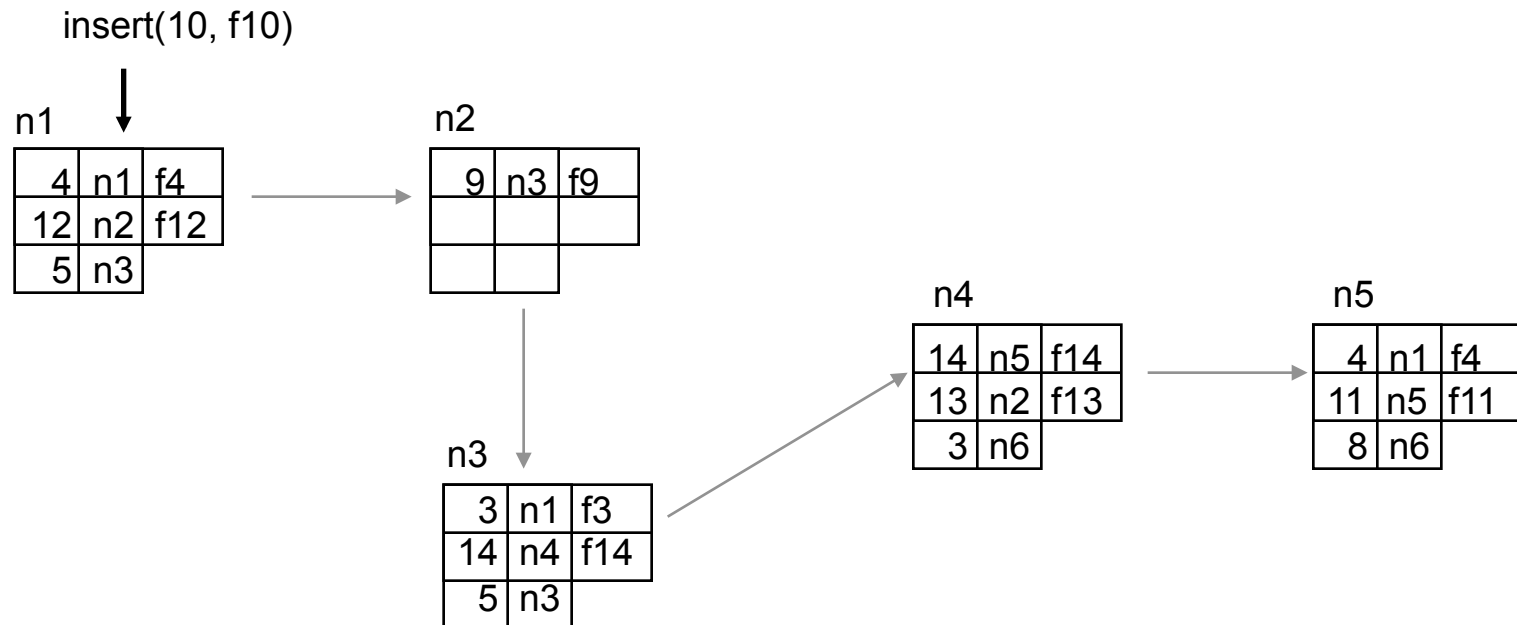
- ❑ When file is returned it is cached along the reverse path (any advantage?)

Insert

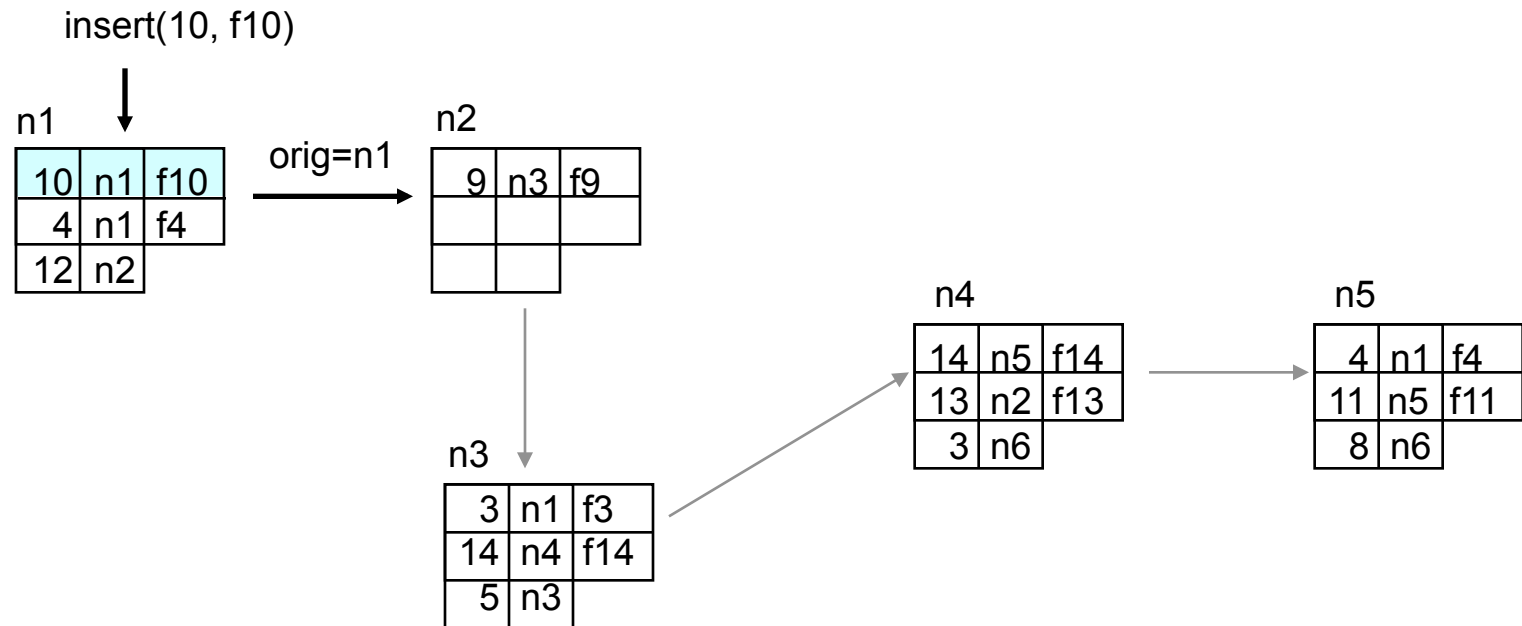
- ❑ API: `insert(id, file);`
- ❑ Two steps
 - First attempt a “search” for the file to be inserted
 - If found, report collision
 - If not found, insert the file by sending it along the query path
 - Inserted files are placed on nodes already possessing files with similar keys
 - A node probabilistically replaces the originator with itself (why?)

Insert Example

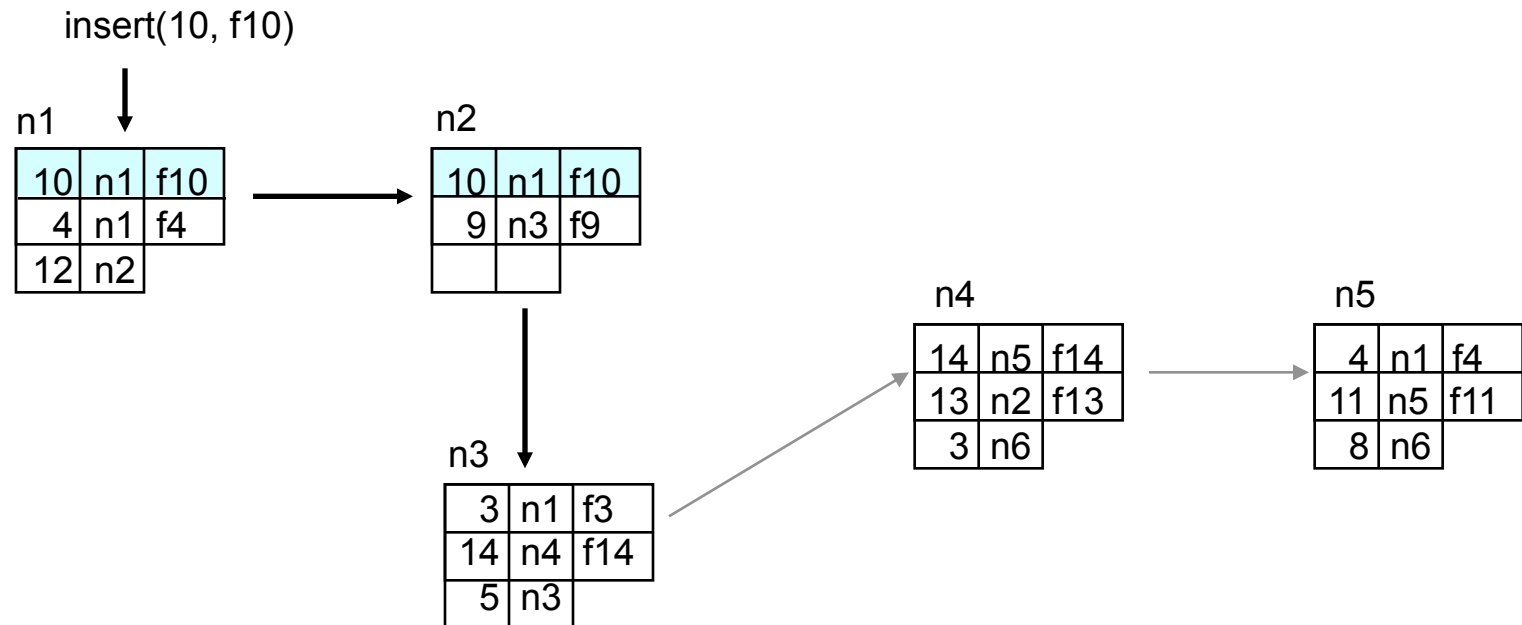
- Assume query returned failure along the shown path (backtrack slightly complicate things); insert f10



Insert Example

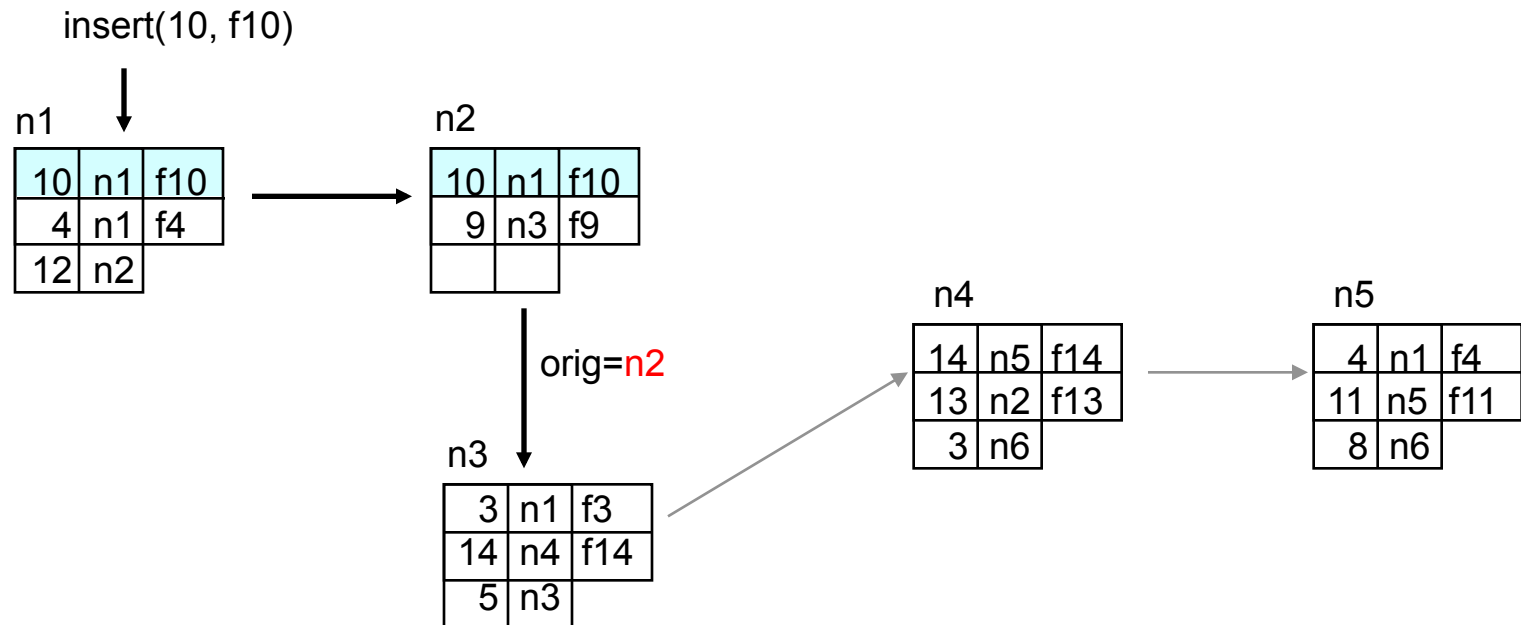


Insert Example




Insert Example


- n2 replaces the originator (n1) with itself



Insert Example



Insert Example




Freenet Analysis

- ❑ Authors claim the following effects:
 - Nodes eventually specialize in locating similar keys
 - If a node is listed in a routing table, it will get queries for related keys
 - Thus will gain “experience” answering those queries
 - Popular data will be transparently replicated and will exist closer to requestors
 - As nodes process queries, connectivity increases
 - Nodes will discover other nodes in the network
- ❑ Caveat: lexicographic closeness of file names/keys may not imply content similarity

Understanding Freenet Self-Organization: Freenet Graph


<i>id</i>	<i>next_hop</i>	<i>file</i>
↓	↓	
	⋮	
↓	↓	
	⋮	

- We create a Freenet reference graph
 - Creating a vertex for each Freenet node
 - Adding a directed link from A to B if A refers to an item stored at B




<i>id</i>	<i>next hop</i>	<i>file</i>
↓	↓	
	⋮	
↓	↓	
	⋮	

- $$1000 / 8 = 125$$



Experiment: Evolution of Freenet Graph



- At each step
 - Pick a node randomly
 - Flip a coin to determine search or insert
 - If search, randomly pick a key in the network
 - If insert, pick a random key




Evolution of path length and clustering;
Clustering is defined as percentage of
local links

Freenet Evolves to Small-World Network

- With usage, the regular, highly localized Freenet network evolved into one irregular graph
- High percentage of highly connected nodes provide shortcuts/bridges
 - Make the world a “small world”
 - Most queries only traverse a small number of hops to find the file



Small-World




❑ First discovered by Milgrom

- In 1967, Milgram mailed 160 letters to a set of randomly chosen people in Omaha, Nebraska
- Goal: pass the letters to a given person in Boston
 - Each person can only pass the letter to an intermediary known on a first-name basis
 - Pick the person who may make the best progress
- Result: 42 letters made it through !
- Median intermediaries was 5.5---thus *six degree of separation*
- A potential explanation: highly connected people with non-local links in mostly locally connected communities improve search performance!
 - 1/4 of all the chains reaching the target person passed through a single person
 - Half the chains were mediated by just three people


Kleinberg's Result on Distributed Search

- ❑ Question: how many long distance links to maintain so that distributed (greedy) search is effective?
- ❑ Assume that the probability of a long link is some (α) inverse-power of the number of lattice steps
- ❑ Kleinberg's Law: Distributed algorithm exists only when probability is proportional to $(\text{lattice steps})^{-d}$, where d is the dimension of the space




Distributed Search

- In other words, if double distance, increase number of neighbors by a constant
-> see Chord



probability is proportional to (lattice steps)^{-d}

Small World



Saul Steinberg; View of World from 9th Ave

Freenet: Properties

- ❑ Query using intelligent routing
 - Decentralized architecture → robust
 - Avoid flooding → low overhead
 - DFS search guided by closeness to target
- ❑ Integration of query and caching makes it
 - Adaptive to usage patterns: reorganize network reference structure
 - Free speech: attempts to discover/supplant existing files will just spread the files !
- ❑ Provide publisher anonymity, security
 - Each node probabilistically replaces originator with itself

Freenet: Issues

- ❑ Does **not** always guarantee that a file is found, even if the file is in the network
- ❑ Good average-case performance, but a potentially **long search path** in a large network
 - Approaching small-world...

Summary

- ❑ All of the previous P2P systems are called unstructured P2P systems
- ❑ Advantages of unstructured P2P
 - Algorithms tend to be simple
 - Can optimize for properties such as locality
- ❑ Disadvantages
 - Hard to make performance guarantee
 - Failure even when files exist